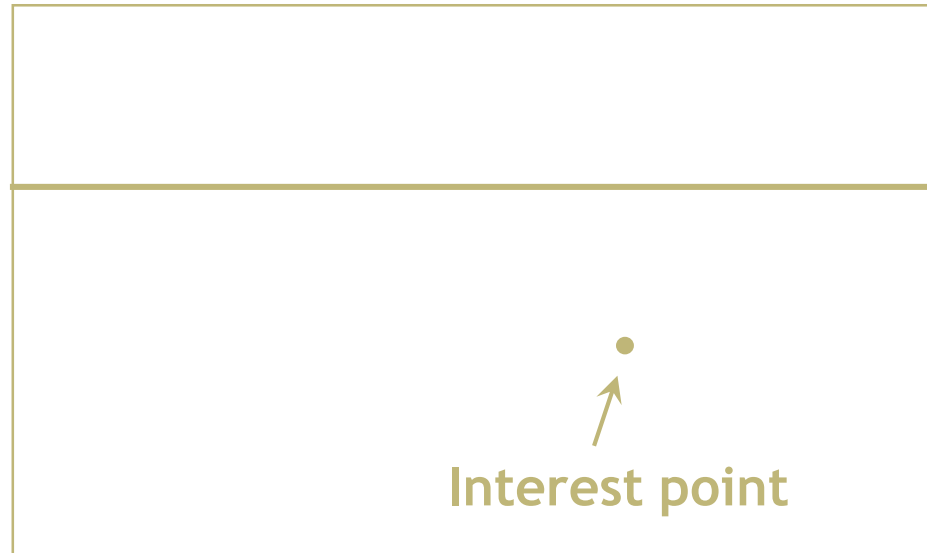


Today- Image Features

-
- Local Features
- Methods for feature extraction

Local Features

- Edges
 - Local in one dimension



Local Features

- Edges
- Corners
- Junctions

- L

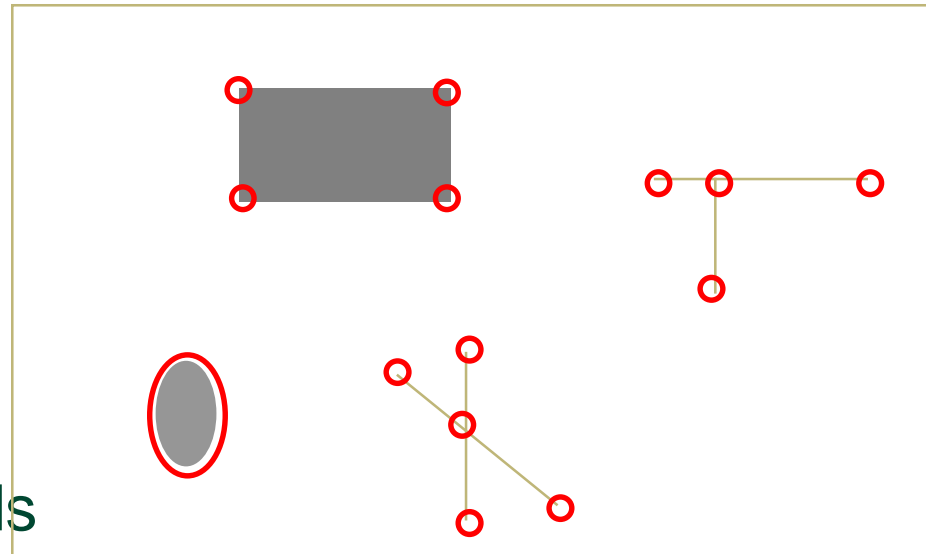
- X

- T

- ...

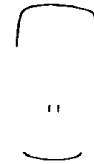
- Edge ends

- Blobs



Local Features

135

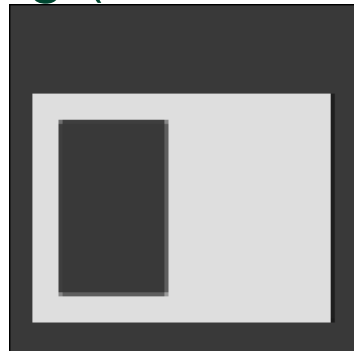


Local Features

- Edge detection
 - One dimensional signal change

Edge Detection

- Principle:
 - find transitions of regions by extracting the edges of regions
 - assumption: regions are (nearly) homogenous
 - physical definition: edges correspond to discontinuities between 'homogeneous' regions
 - we have to take into account the noise
 - edge detection using (1st or 2nd) derivatives



Edge Detection

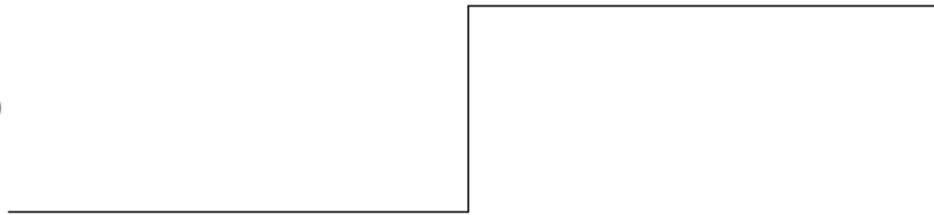
- What are edges ?
 - Humans find object boundaries very quickly - which is why we would like to define 'edges' to be the object boundaries
 - to detect object boundaries we have to know which objects are in the scene (I.e. we have to recognize the object(s) first...):



What are 'edges' (1D)

- Idealized:

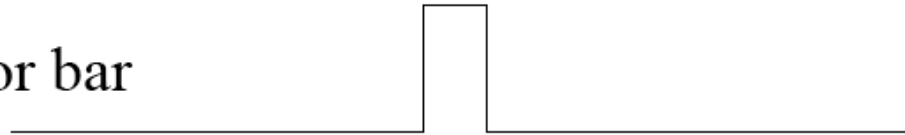
step



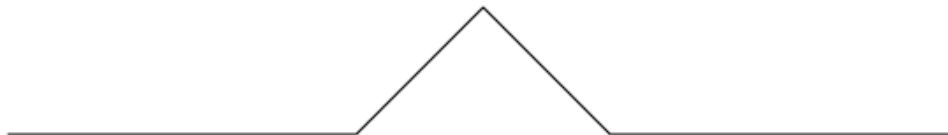
ramp



line or bar



roof



Edge Detection

- What are edges ?

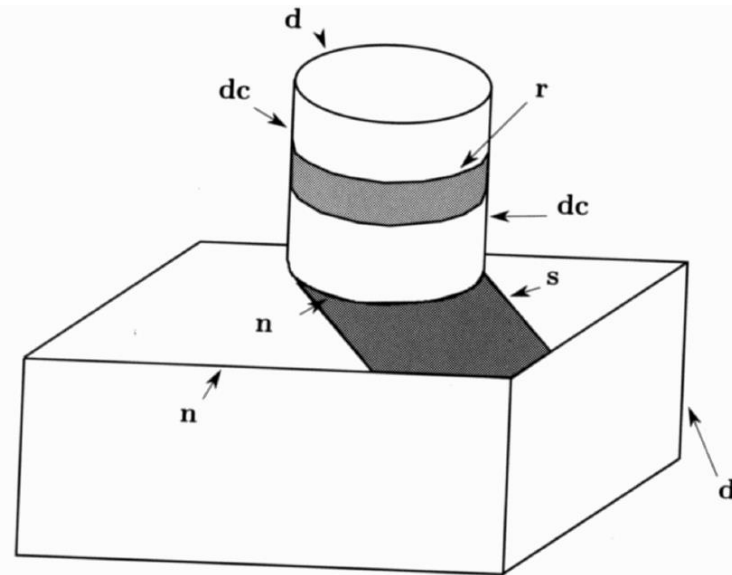
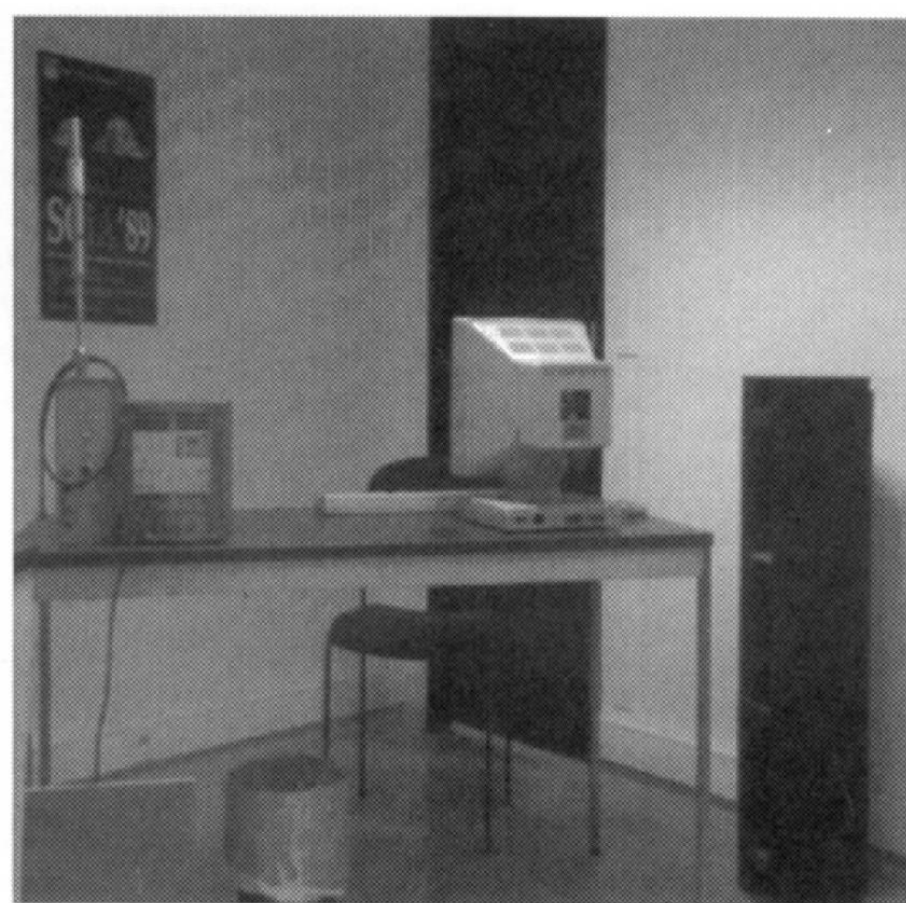


Figure 4.1 Edges in an image have different physical sources.

- object-background boundaries (edges (d) and depth discontinuity (dc))
- object-object boundaries (those correspond not necessarily to object boundaries, (n))
- shadows (s)
- discontinuities of object texture (r)
- discontinuities of surface normals (n)

Example

- Image with multiple objects:



2D Edge Detection

- calculate derivative
 - use the magnitude of the gradient
 - the gradient is:

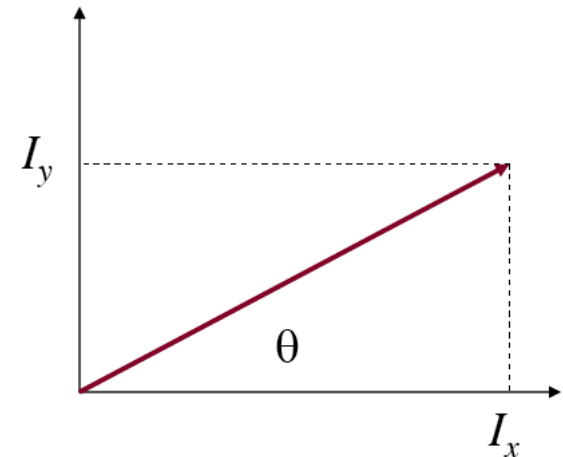
$$\nabla I = (I_x, I_y) = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

- the magnitude of the gradient is:

$$\|\nabla I\| = \sqrt{I_x^2 + I_y^2}$$

- the direction of the gradient is:

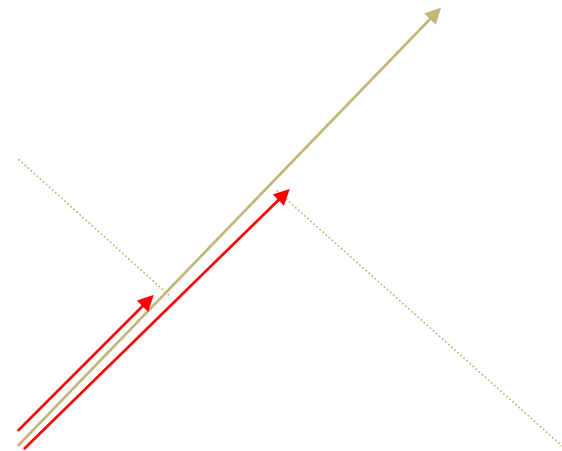
$$\theta = \arctan(I_y, I_x)$$



'Steerable' Filters

- what is the gradient in some direction θ ?

$$f'_\theta(x, y) = \frac{\partial f}{\partial x} \cos \theta + \frac{\partial f}{\partial y} \sin \theta$$



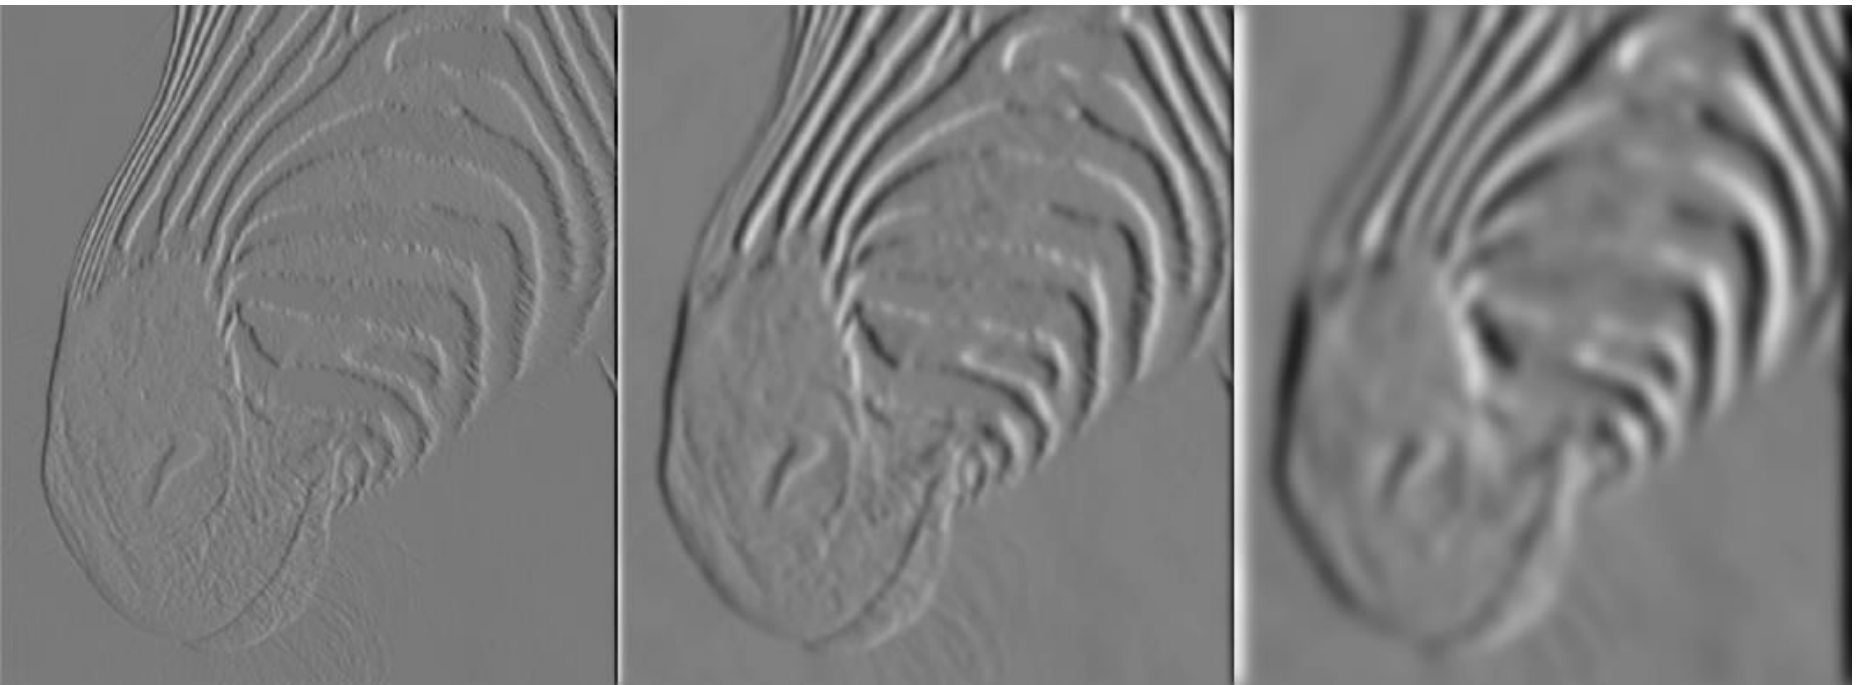
2D Edge Detection

- the scale of the smoothing filter affects derivative estimates, and also the semantics of the edges recovered
 - note: strong edges persist across scales

1 pixel

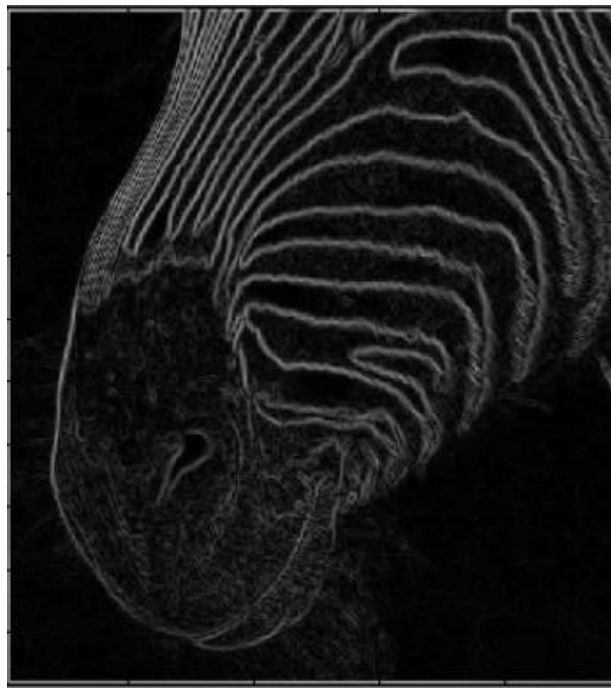
3 pixels

7 pixels



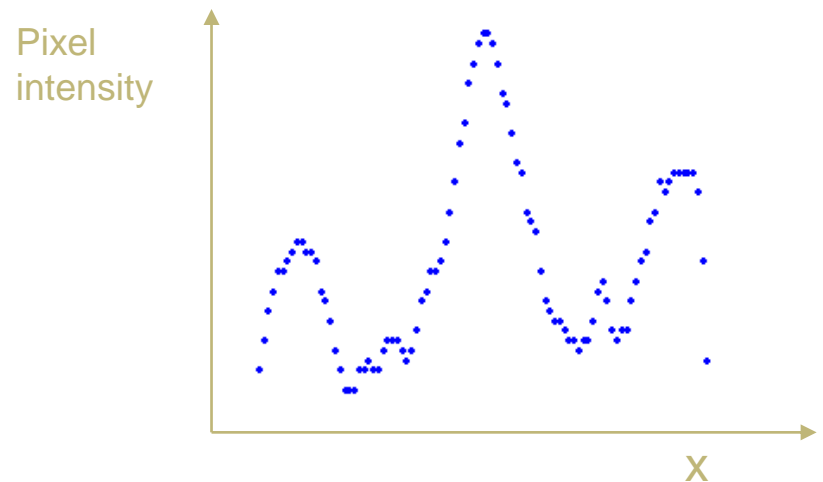
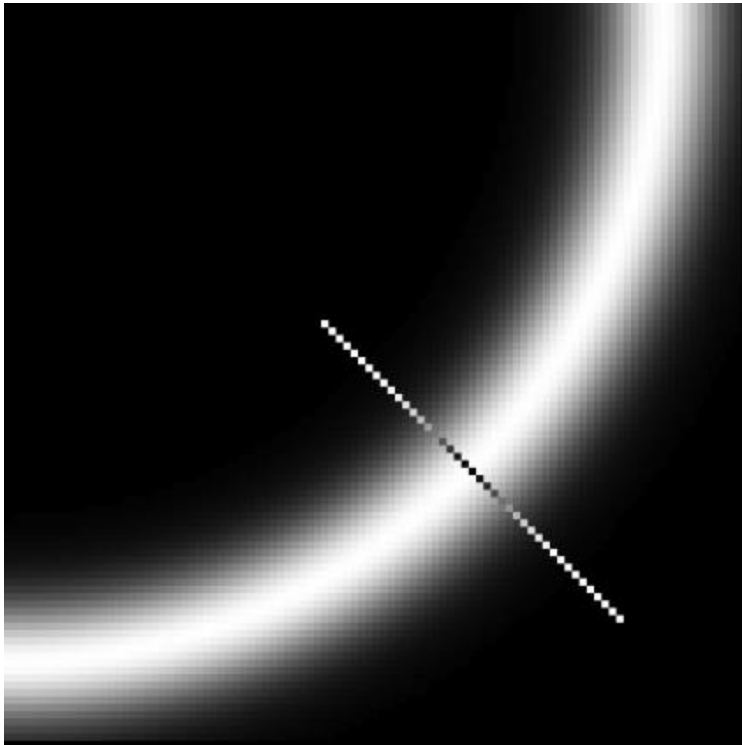
2D Canny Edge Detection

- there are 3 major issues:
 - the gradient magnitude at different scales is different; which should we choose?
 - the gradient magnitude is large along a thick trail; how do we identify the significant points?
 - how do we link the relevant points up into curves?



2D Canny Edge Detection

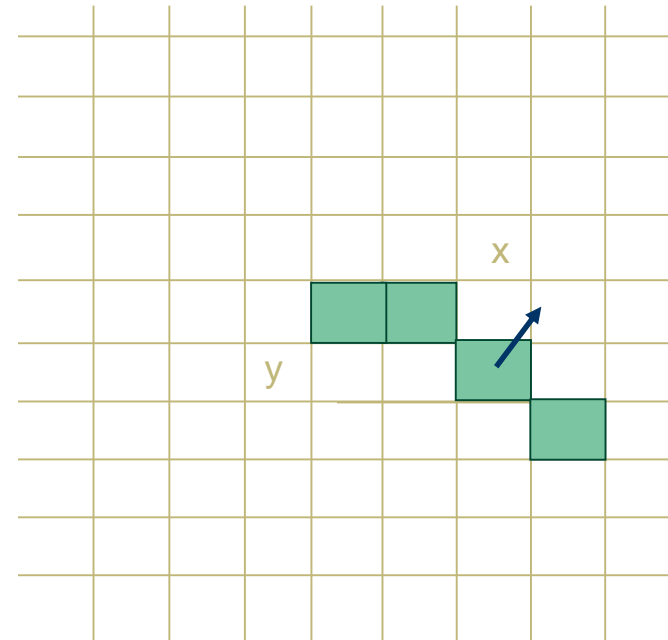
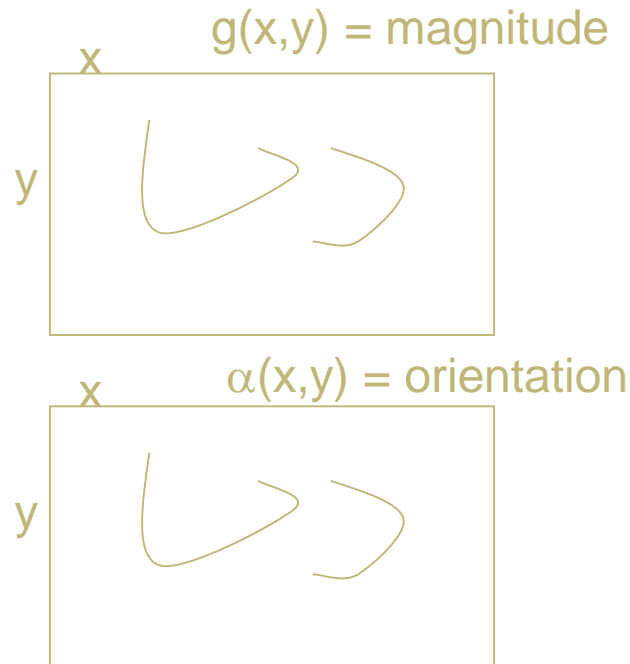
- Non-Maxima Suppression
- look in a neighborhood along the direction of the gradient
- choose the largest gradient magnitude in this neighborhood



If $f(x) > f(x-1) \ \& \ f(x) > f(x+1)$ $\Rightarrow \ x = 1$
else $x = 0$

2D Edge Detection

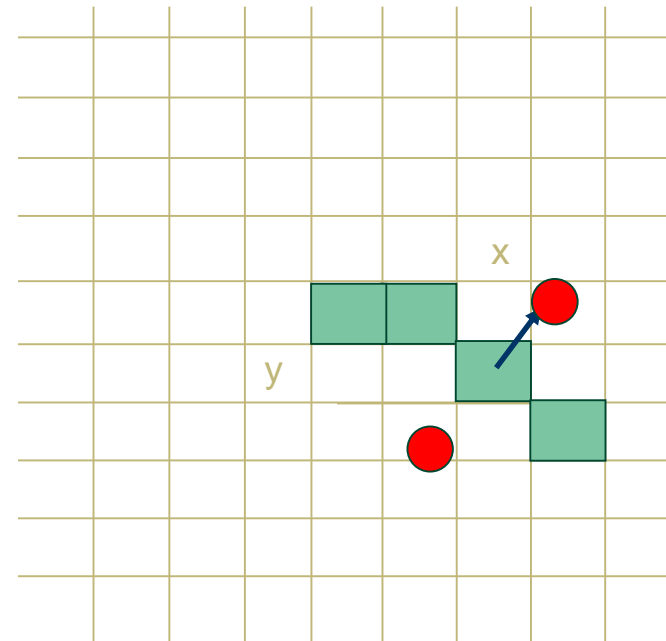
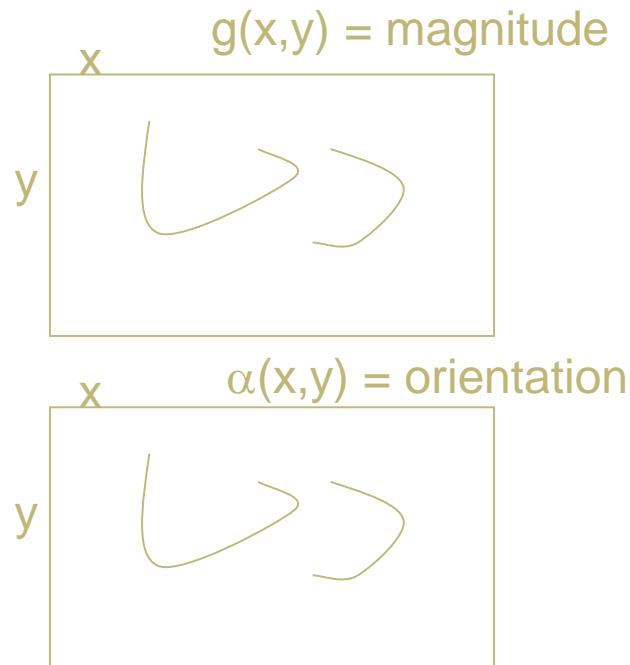
- Non-Maxima Suppression



Gradient magnitude $g(x,y)$, orientation $\alpha(x,y)$

2D Edge Detection

- Non-Maxima Suppression



Gradient magnitude $g(x,y)$, orientation $\alpha(x,y)$

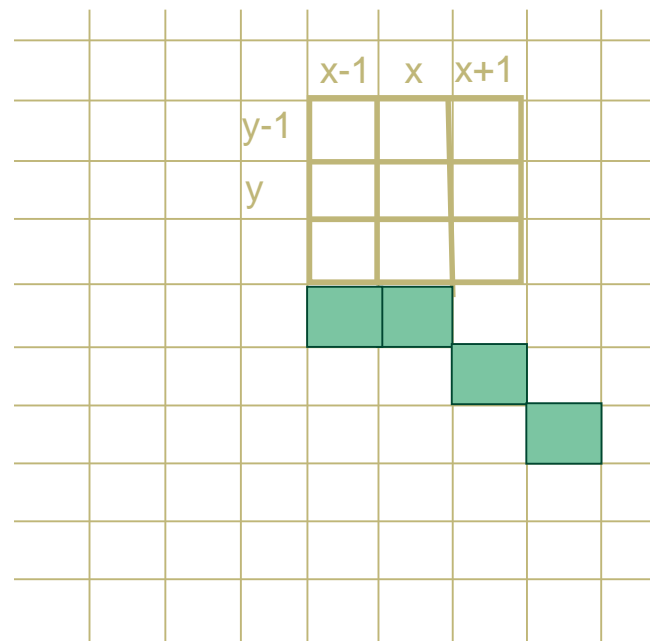
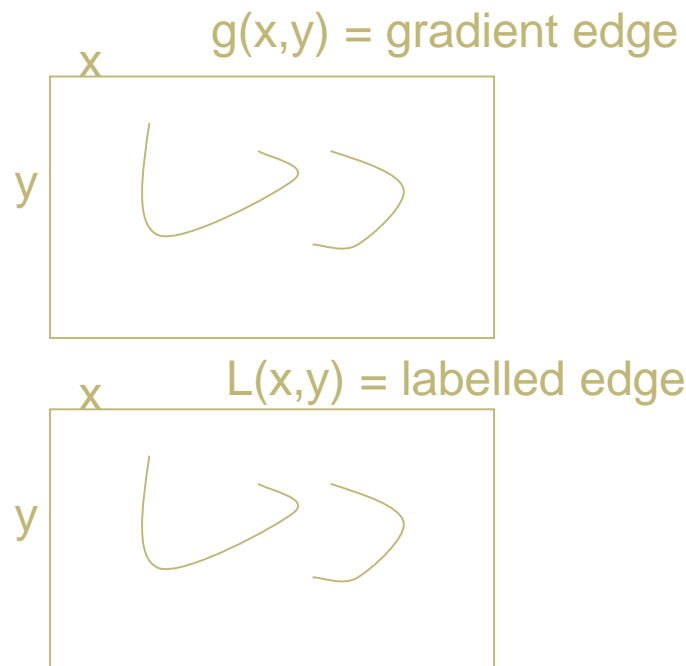
If $g(x,y) > g(x+dx,y+dy)$ && If $g(x,y) > g(x-dx,y-dy)$ \leq bilinear interpolation

else $g(x,y) = 0$

Where $\alpha(x,y) = \text{atan}(dy,dx)$

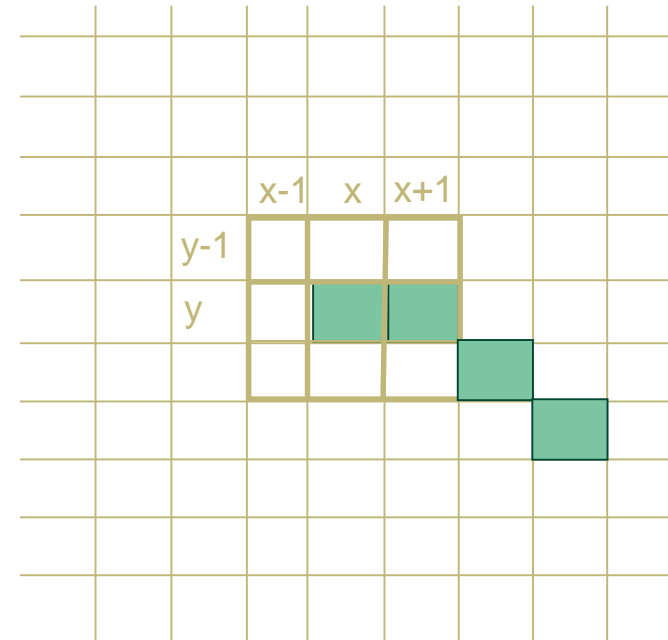
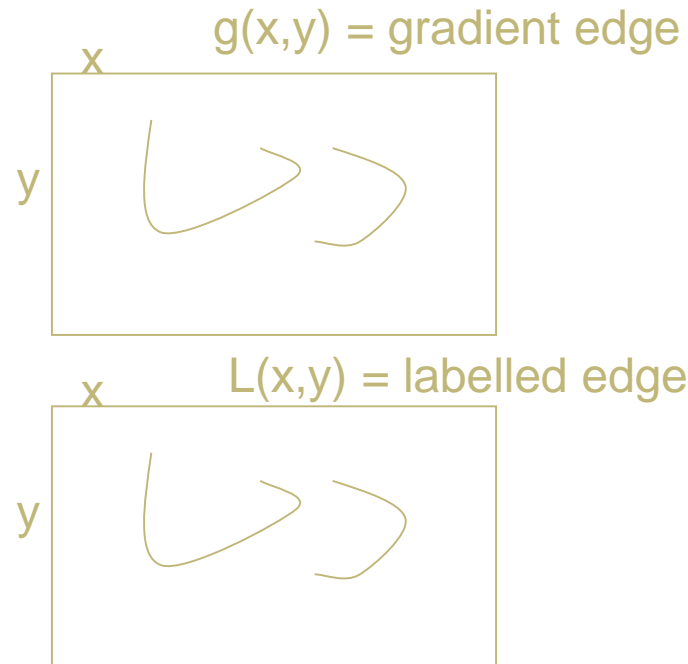
2D Edge Detection

- Labelling connected edges



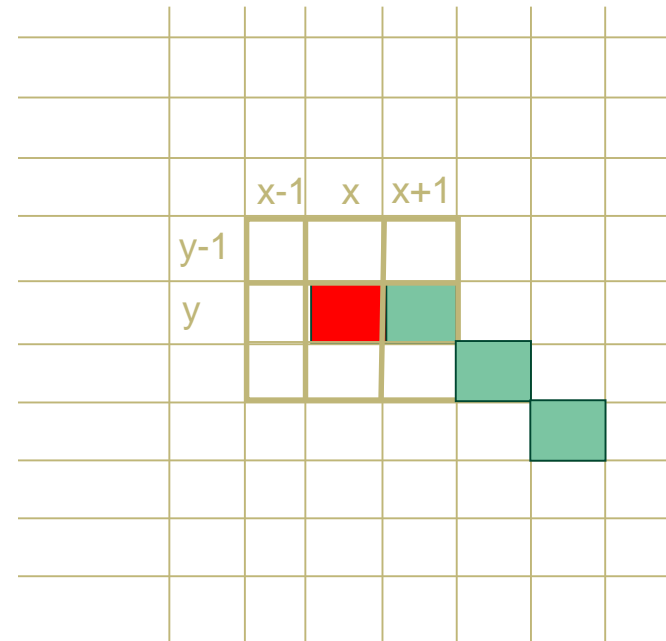
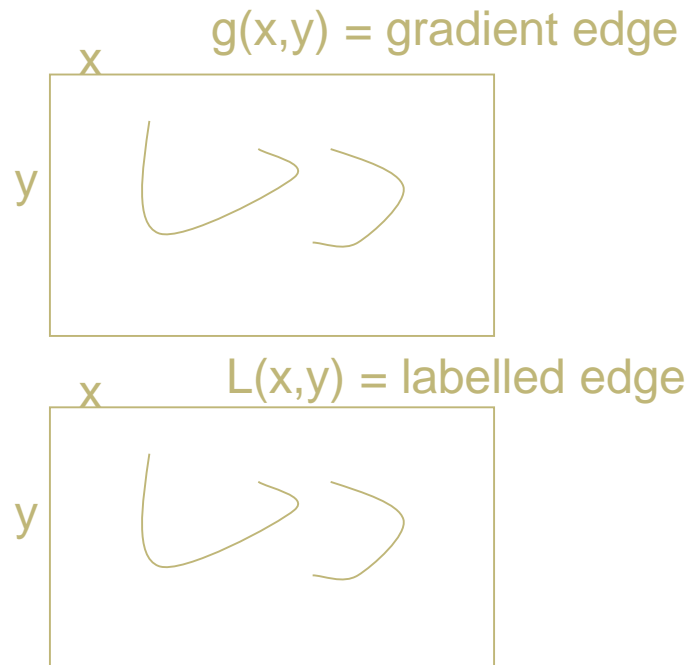
2D Edge Detection

- Labelling connected edges



2D Edge Detection

- Labelling connected edges

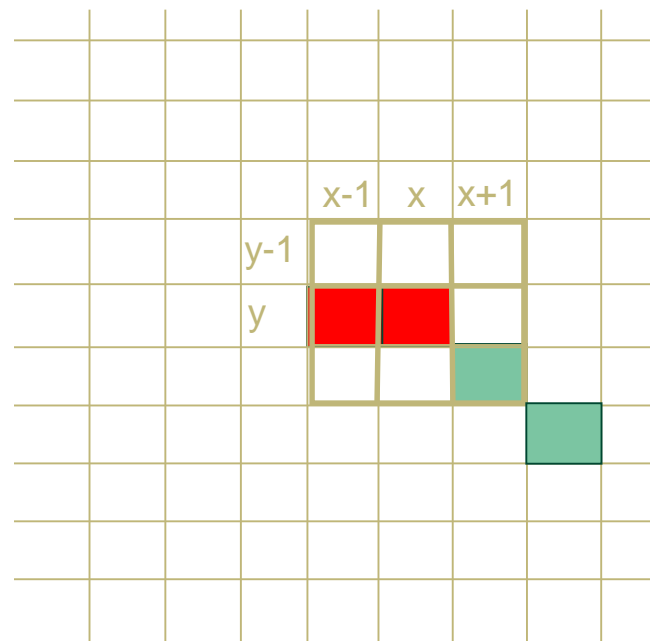
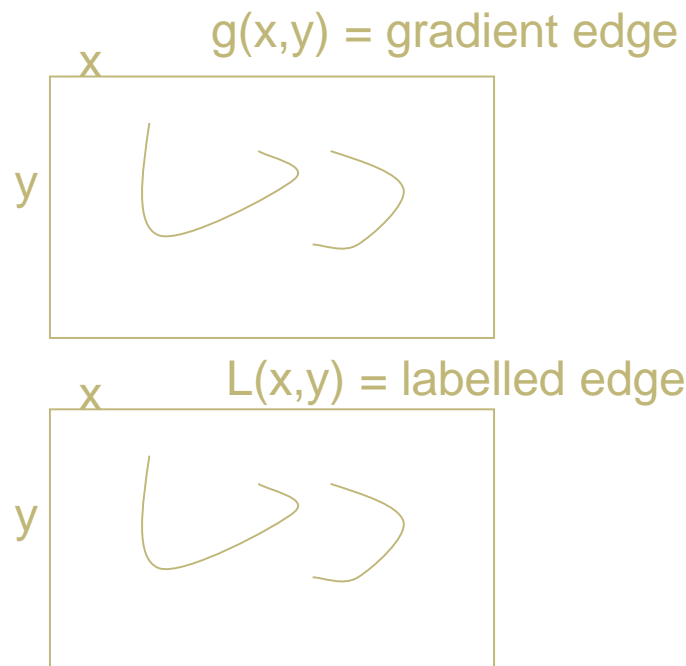


If $g(x,y) > TH$ $L(x,y) = \{L(x-1,y-1) \parallel L(x,y-1) \parallel L(x+1,y-1) \parallel L(x-1,y) \parallel \text{new_label}\}$

else If $g(x,y) > TL$ $L(x,y) = t_edge$

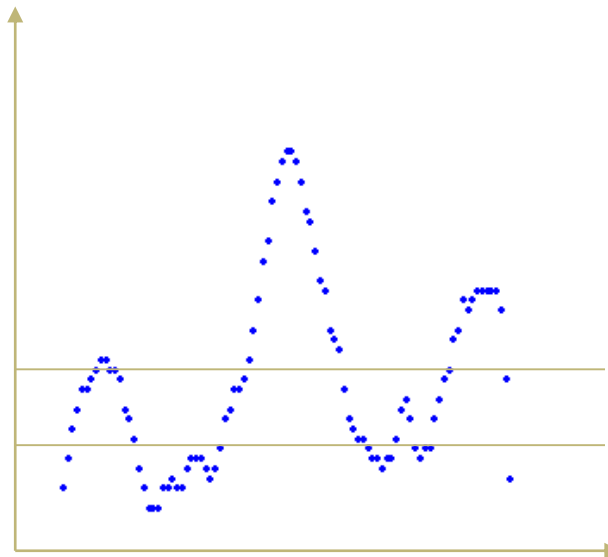
2D Edge Detection

- Labelling connected edges



2D Edge Detection

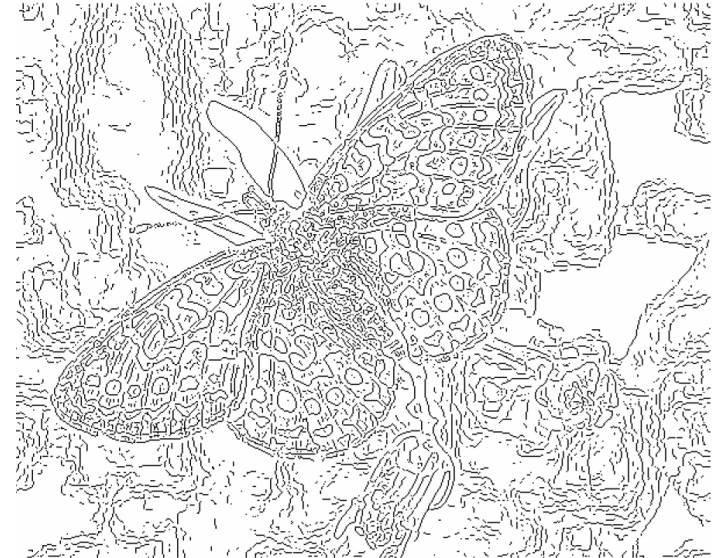
- Hysteresis Thresholding
 - High and low Threshold
 - High threshold to validate the edge points
 - Low threshold to remove the noise



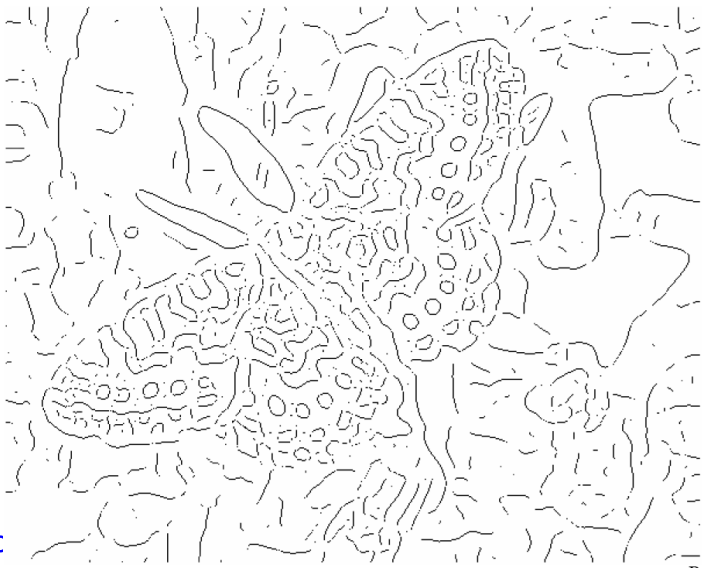
Canny Edge detection

- Computation of image derivatives
- Conversion to gradient magnitude and orientation
- Detection of local maxima in the gradient direction (non maxima suppression)
- Labelling connected edges
- Hysteresis thresholding

Butterfly Example



fine scale
high
threshold



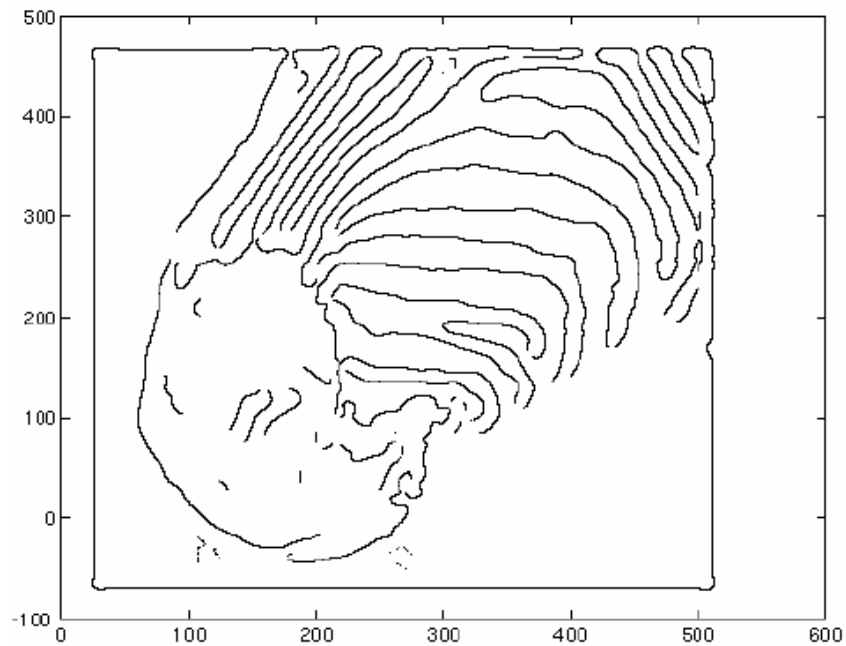
coarse
scale
low
threshold



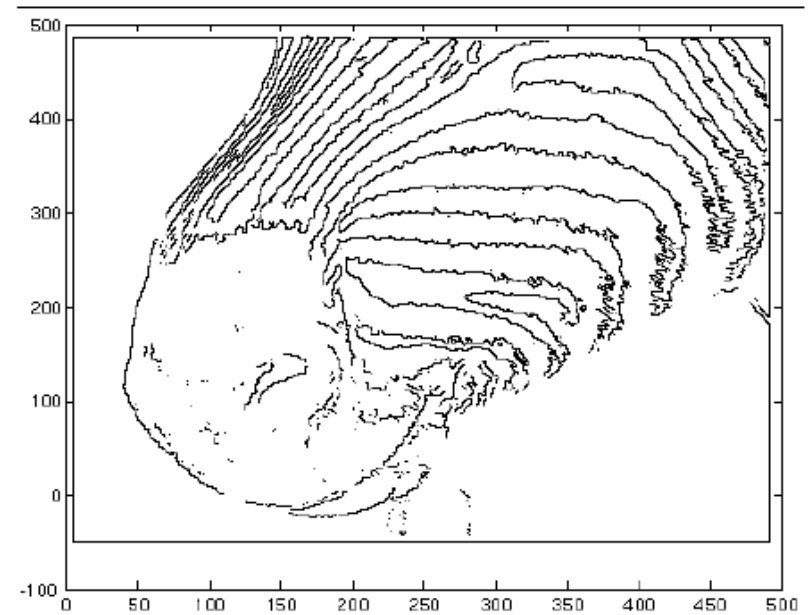
coarse
scale,
high
threshold

Edges...

- $\sigma = 4$

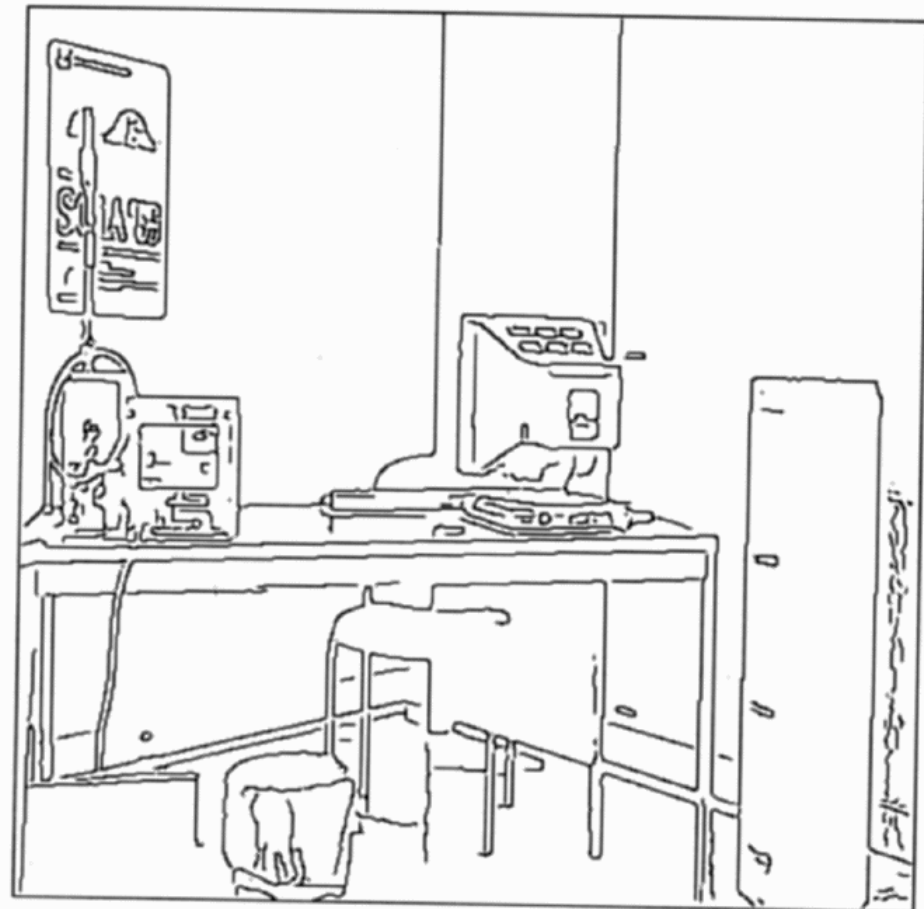
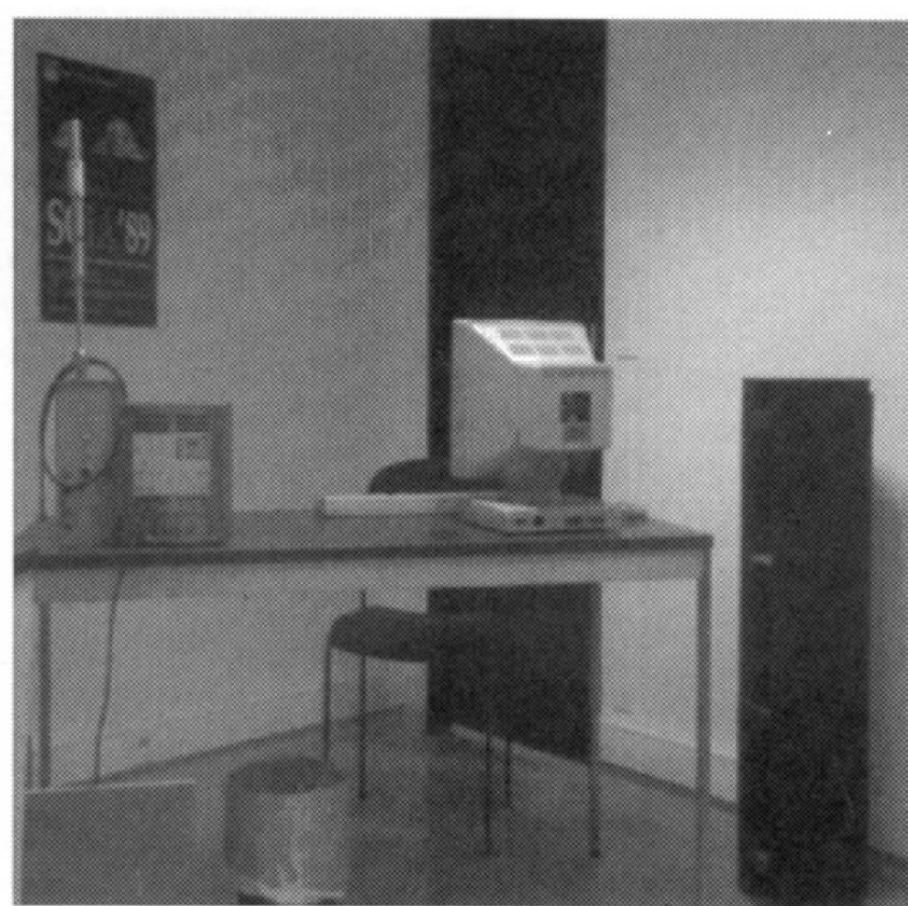


- $\sigma = 2$



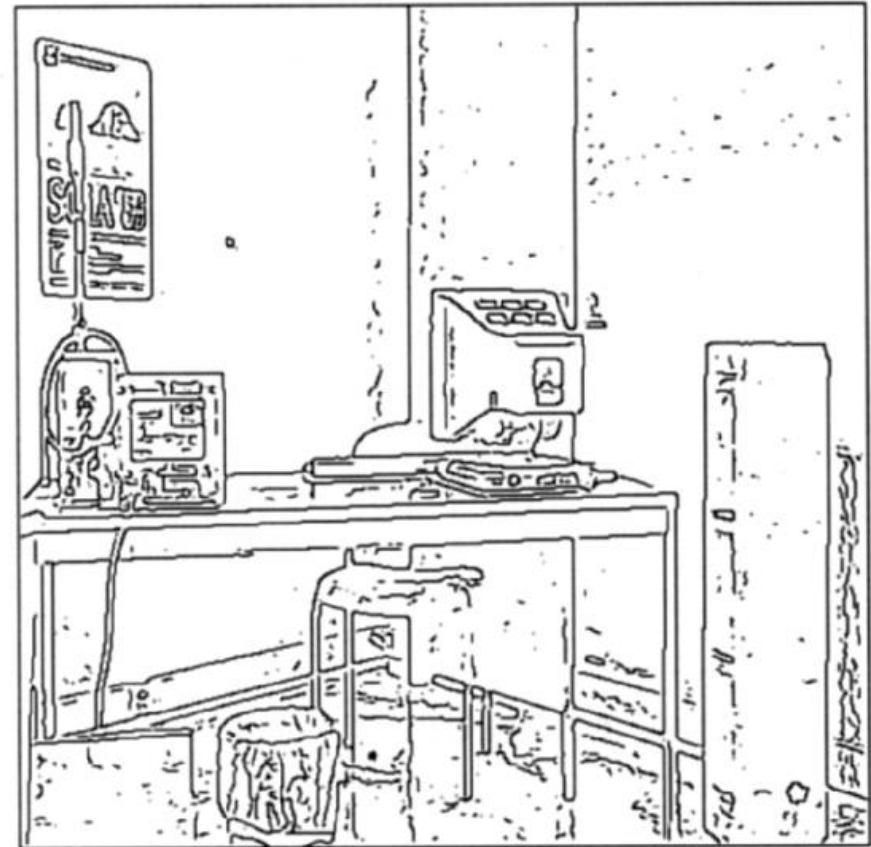
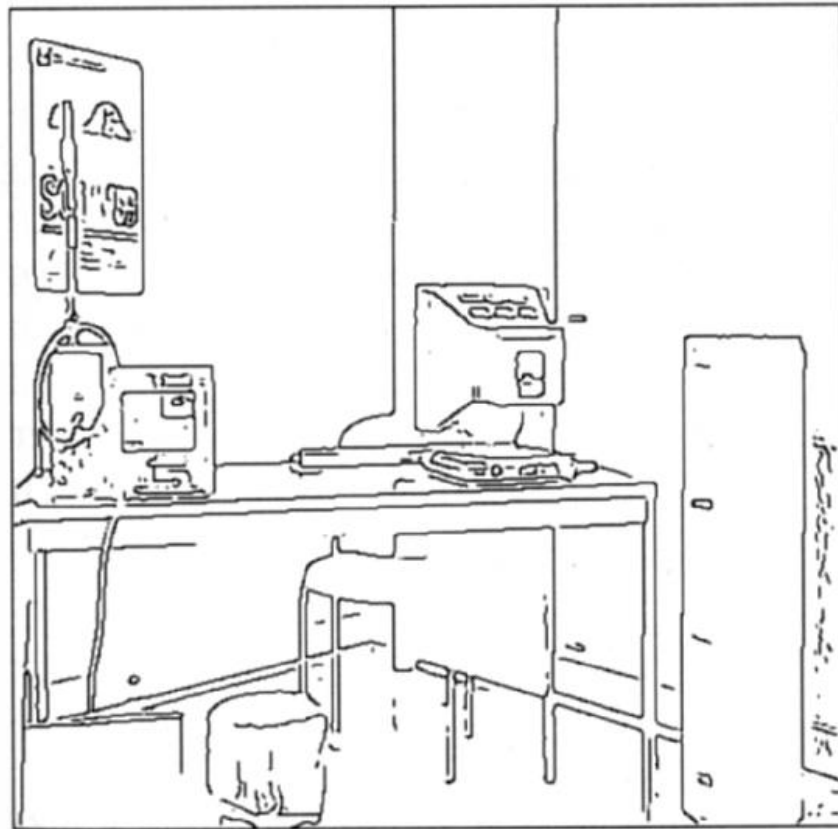
Edges...

- contour following with 2 thresholds

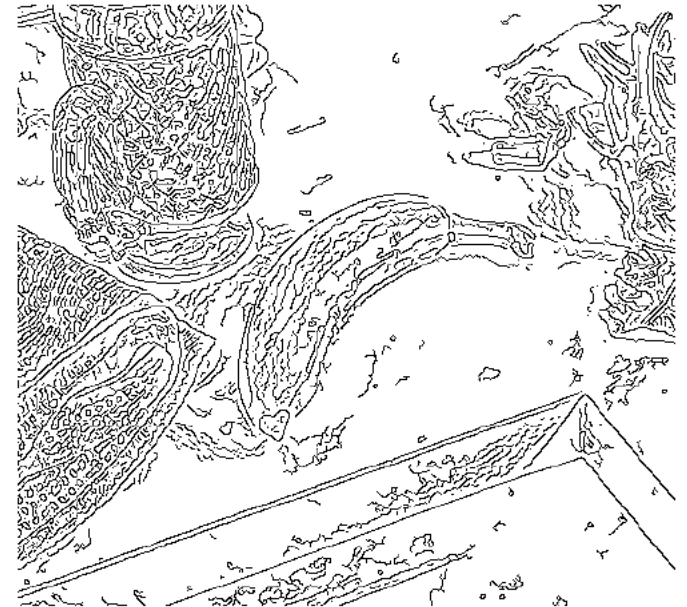


Edges...

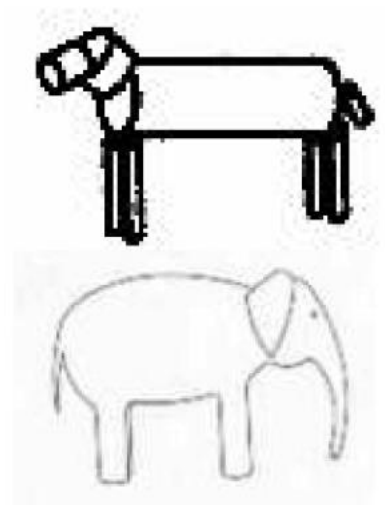
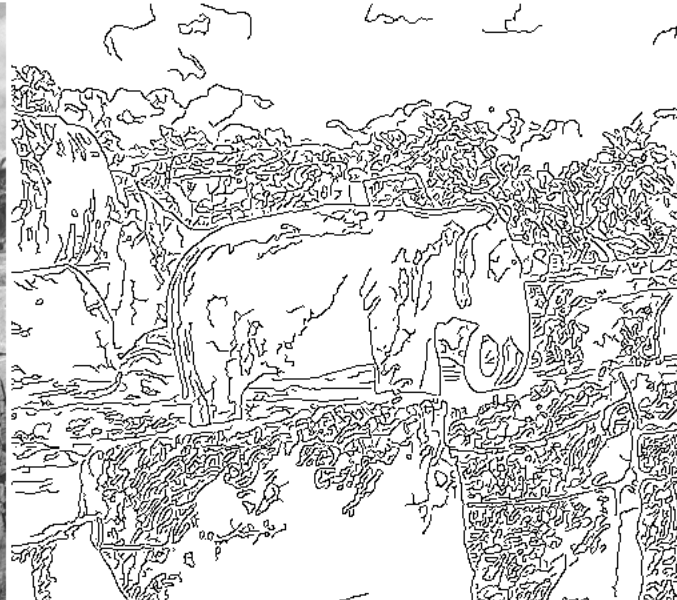
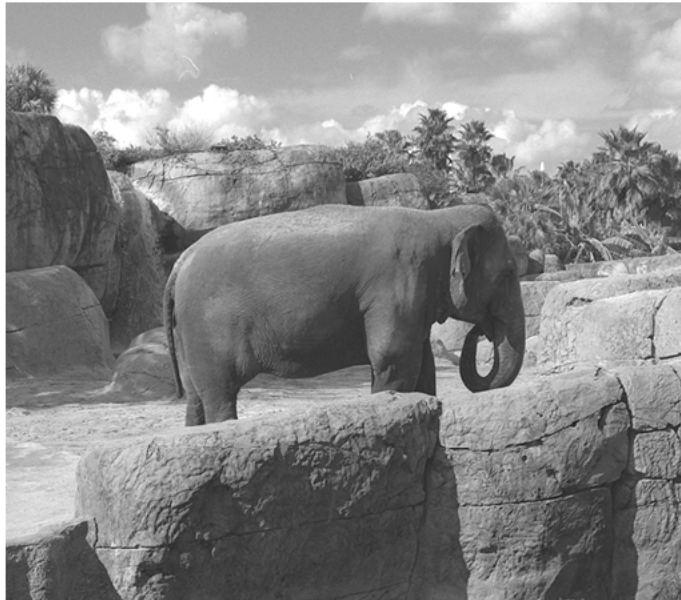
- different thresholds



line drawing vs. edge detection



Edges...



Match “model” to
measurements?

Local Features

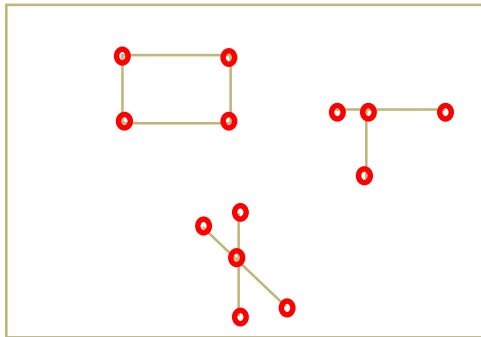
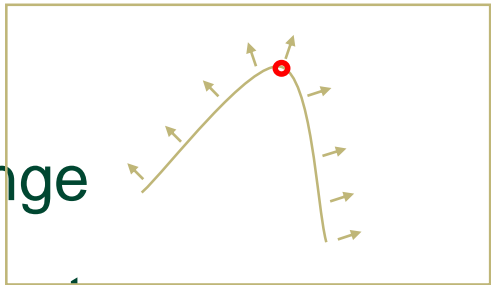
- Corners, blobs
 - Two dimensional signal change
 - More complex local structures



Feature detectors

Contour based methods

- Detecting curvature change
 - Detecting edges
 - Detecting sudden edge orientation change
- Detecting intersections of line segments
 - Detecting edges
 - Fitting line segments to the edges
i.e., Hough transform
 - Finding intersections



Feature detectors

- Intensity based methods

Eigenvalues-reminder

- Singular value decomposition

$$A = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} = \begin{bmatrix} u_{11} & u_{12} \\ u_{21} & u_{22} \end{bmatrix} \cdot \begin{bmatrix} d_1 & 0 \\ 0 & d_2 \end{bmatrix} \cdot \begin{bmatrix} v_{11} & v_{12} \\ v_{21} & v_{22} \end{bmatrix}^T = U \cdot D \cdot V^T$$

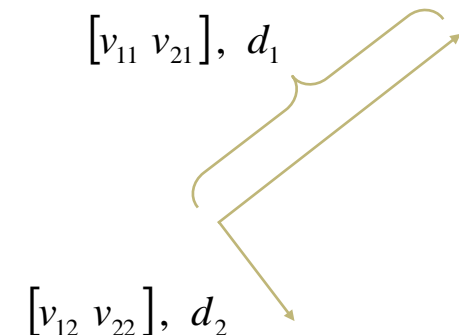
eigenvectors $\begin{bmatrix} v_{11} & v_{21} \end{bmatrix}^T, \begin{bmatrix} v_{12} & v_{22} \end{bmatrix}^T$

$$UU^T = VV^T = 1 \quad U^{-1} = U^T \quad V^{-1} = V^T$$

eigenvalues $d_1, d_2 \geq 0$

Eigenvector, eigenvalue

determinant $\det(A) = ad - cb = d_1 d_2$



Feature detectors

Intensity based methods [Beaudet'78]

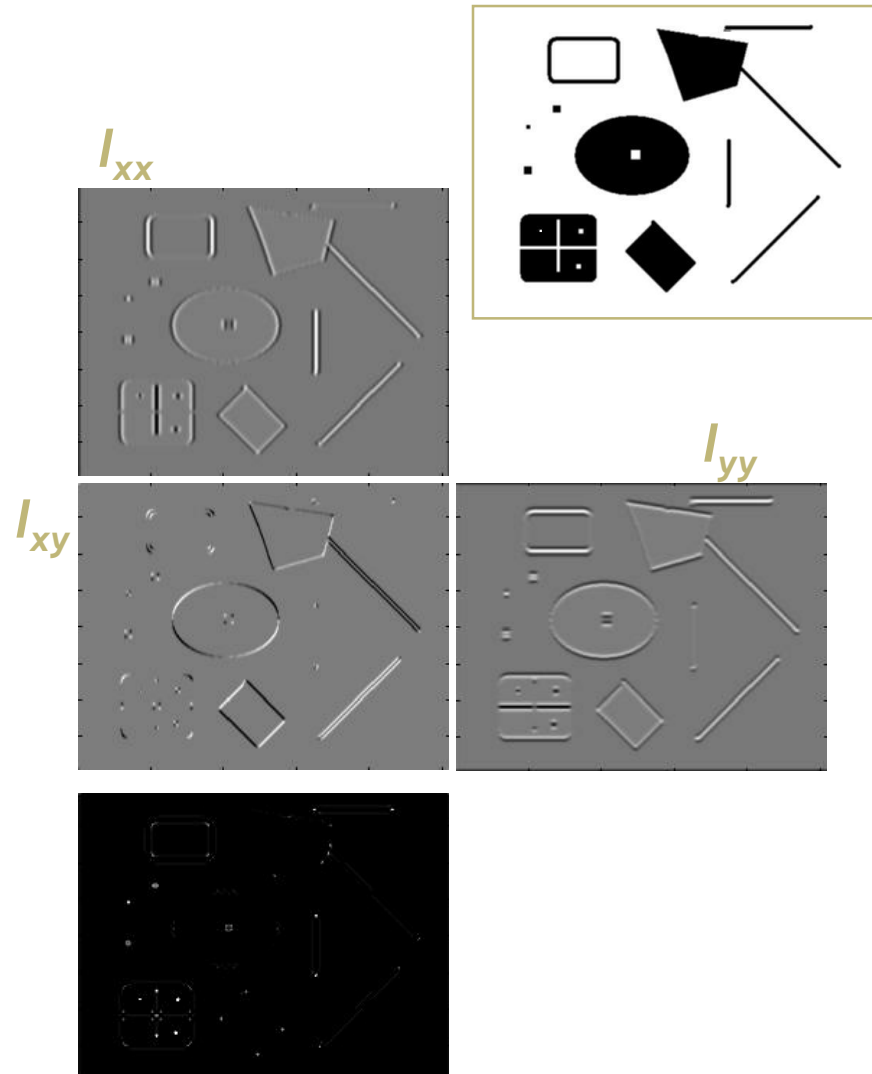
- Hessian determinant

$$Hessian(I) = \begin{bmatrix} I_{xx} & I_{xy} \\ I_{xy} & I_{yy} \end{bmatrix}$$

$$\det(Hessian(I)) = I_{xx}I_{yy} - I_{xy}^2$$

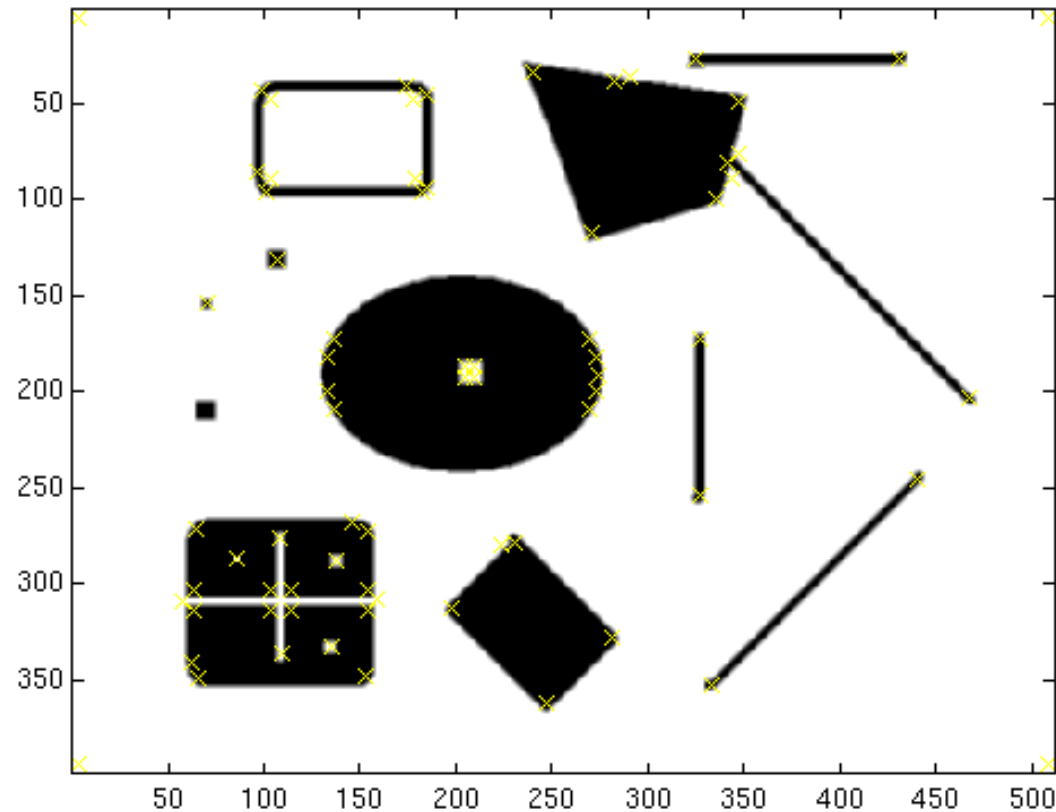
In Matlab:

$$I_{xx}.*I_{yy} - (I_{xy})^2$$



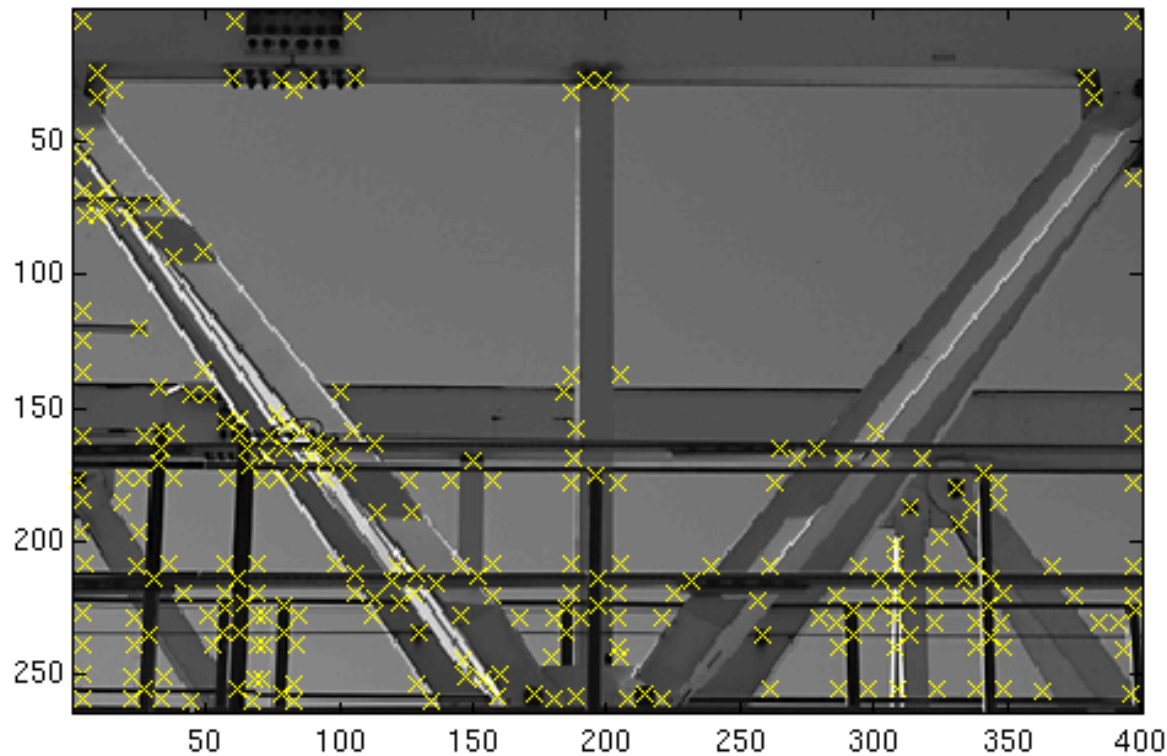
Feature detectors

Intensity based methods [Beaudet'78]



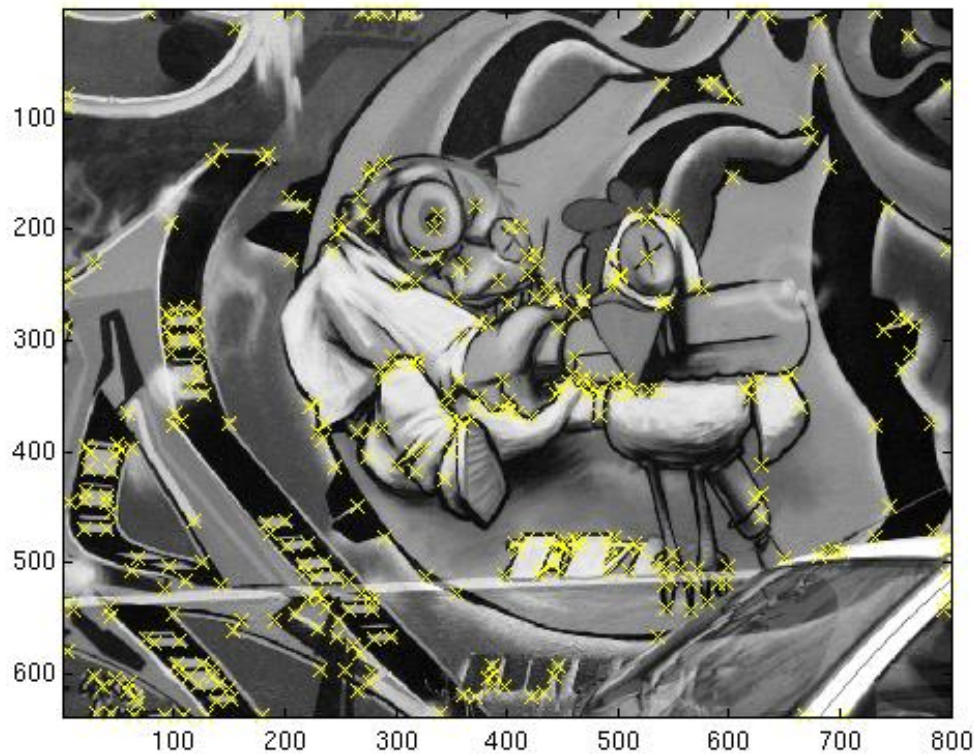
Feature detectors

Intensity based methods [Beaudet'78]



Feature detectors

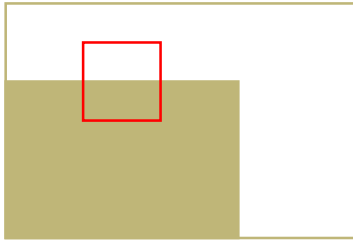
Intensity based methods [Beaudet'78]



Feature detectors

Intensity based methods [Moravec'77]

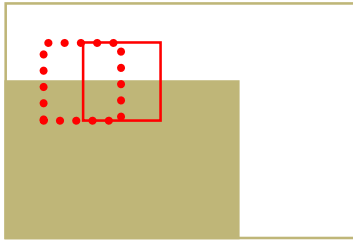
- Autocorrelation function



Feature detectors

Intensity based methods [Moravec'77]

- Autocorrelation function

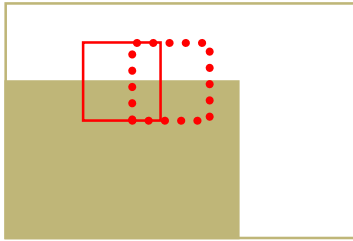


$$(a = \text{[dashed box]} - \text{[solid box]})$$

Feature detectors

Intensity based methods [Moravec'77]

- Autocorrelation function

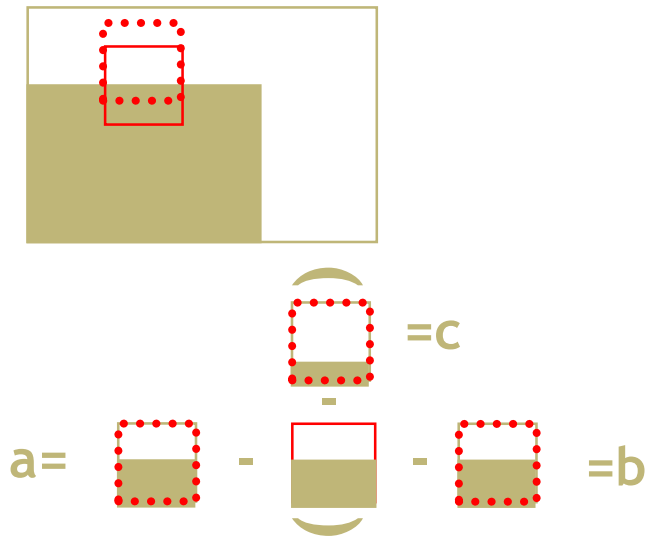


$$a = \left(\text{dotted box} - \left(\text{solid box} - \text{dotted box} \right) \right) = b$$

Feature detectors

Intensity based methods [Moravec'77]

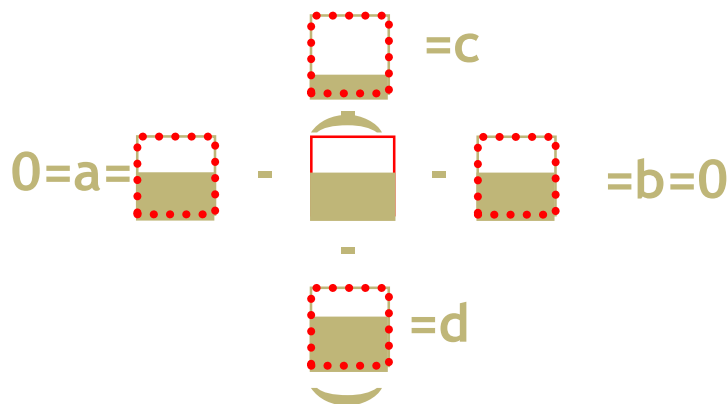
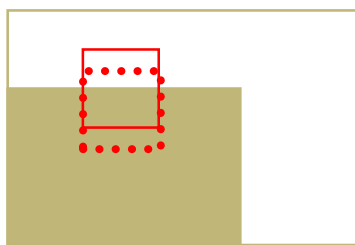
- Autocorrelation function



Feature detectors

Intensity based methods [Moravec'77]

- Autocorrelation function

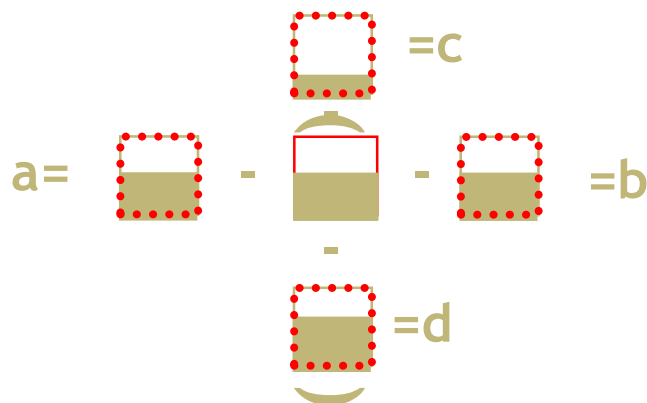
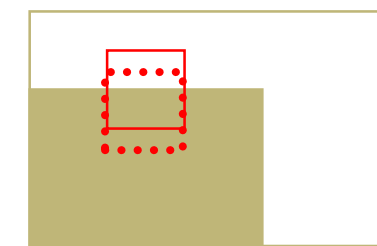


$$\min(a,b,c,d) > T$$

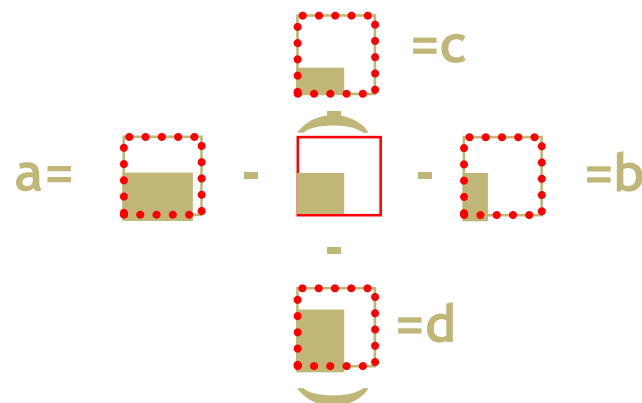
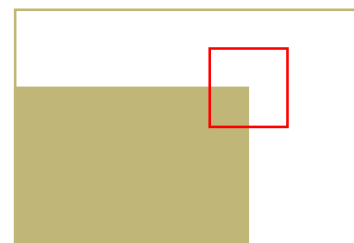
Feature detectors

Intensity based methods [Moravec'77]

- Autocorrelation function



$$\min(a, b, c, d) > T$$



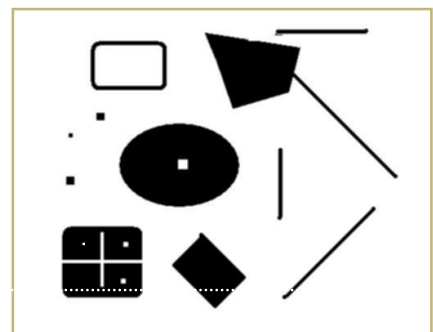
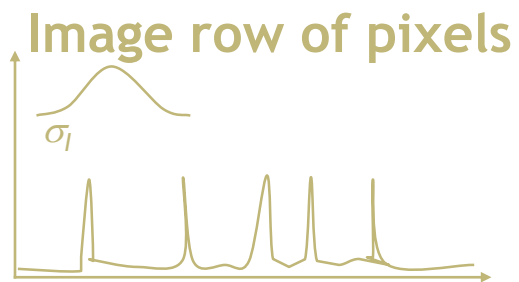
$$\min(a, b, c, d) > T$$

Feature detectors

Intensity based methods [Harris'88]

- Second moment matrix autocorrelation matrix

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$



I_y

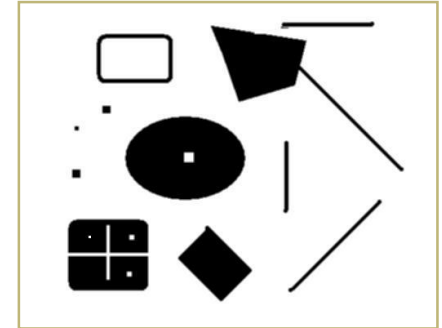
Feature detectors

Intensity based methods [Harris'88]

- Second moment matrix autocorrelation matrix

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

- Image derivatives
 $g_x(\sigma_D)$, $g_y(\sigma_D)$,



Feature detectors

Intensity based methods [Harris'88]

- Second moment matrix autocorrelation matrix

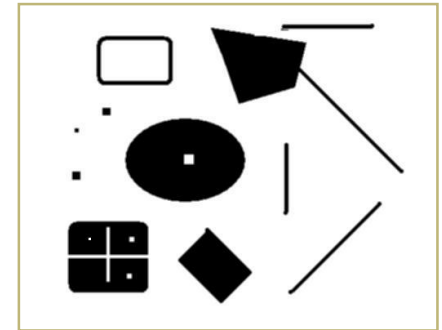
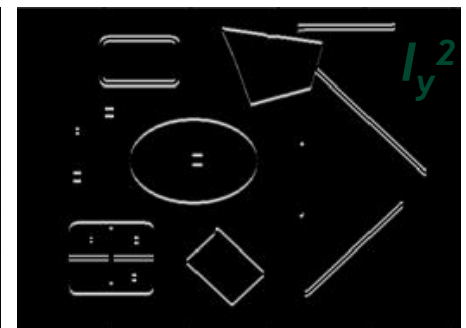
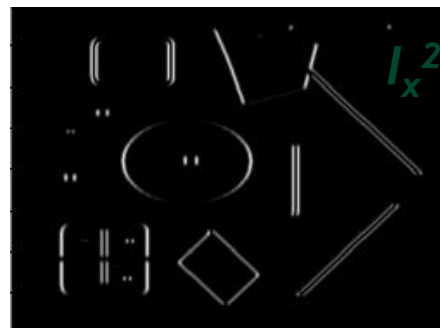
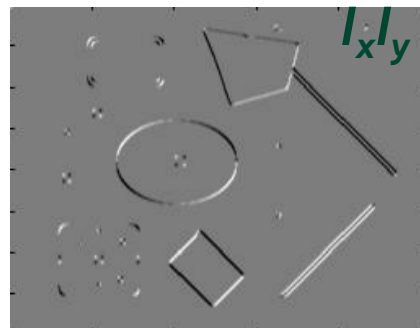
$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives

$$g_x(\sigma_D), g_y(\sigma_D),$$



2. Square of derivatives

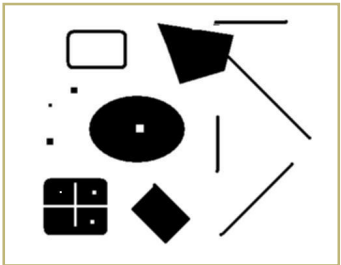


Feature detectors

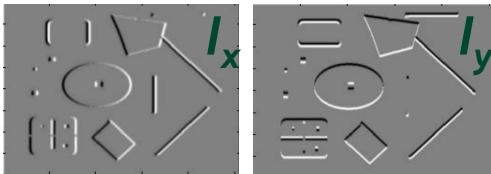
Intensity based methods [Harris'88]

- Second moment matrix autocorrelation matrix

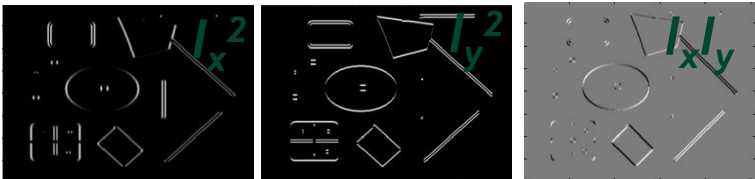
$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$



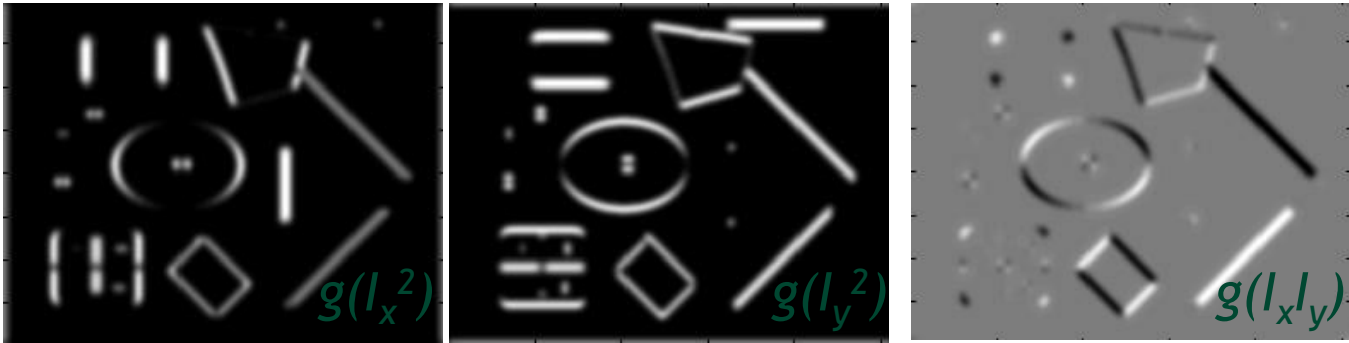
1. Image derivatives



2. Square of derivatives

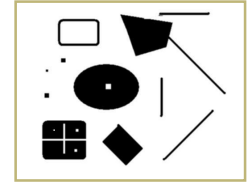


3. Gaussian filter $g(\sigma_I)$



Feature detectors

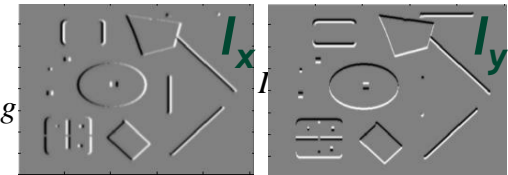
Intensity based methods [Harris'88]



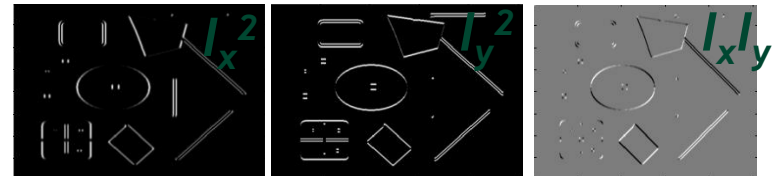
- Second moment matrix autocorrelation matrix

1. Image derivatives

$$\mu(\sigma_I, \sigma_D) = g$$



2. Square of derivatives



3. Gaussian filter $g(\sigma_I)$



4. Cornerness function - both eigenvalues are strong

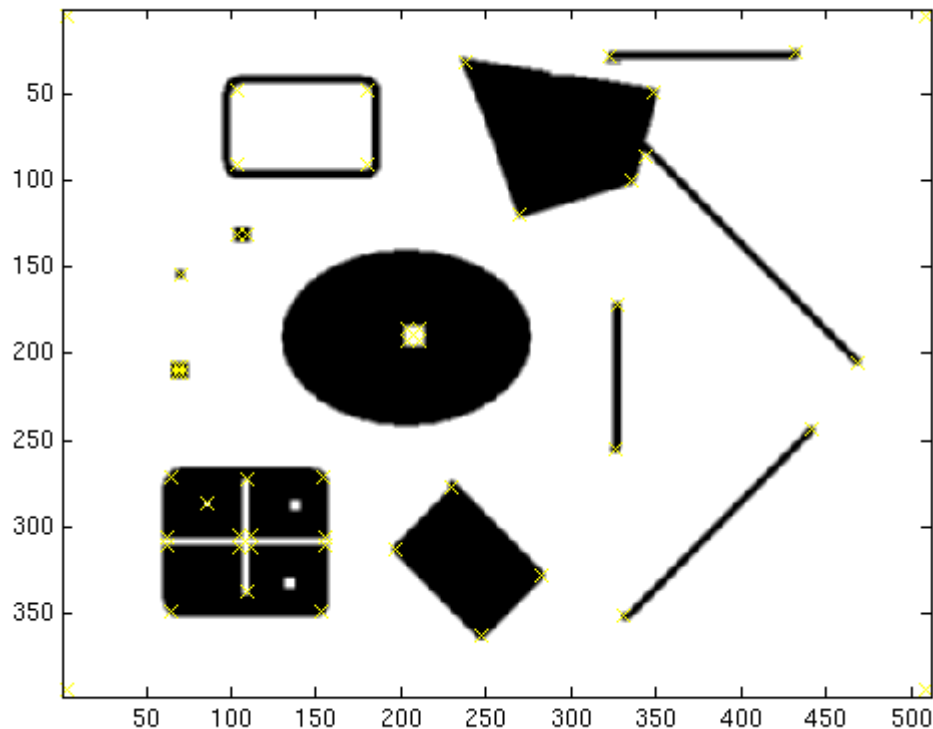
$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha[\text{trace}(\mu(\sigma_I, \sigma_D))] = g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2$$

5. Non-maxima suppression



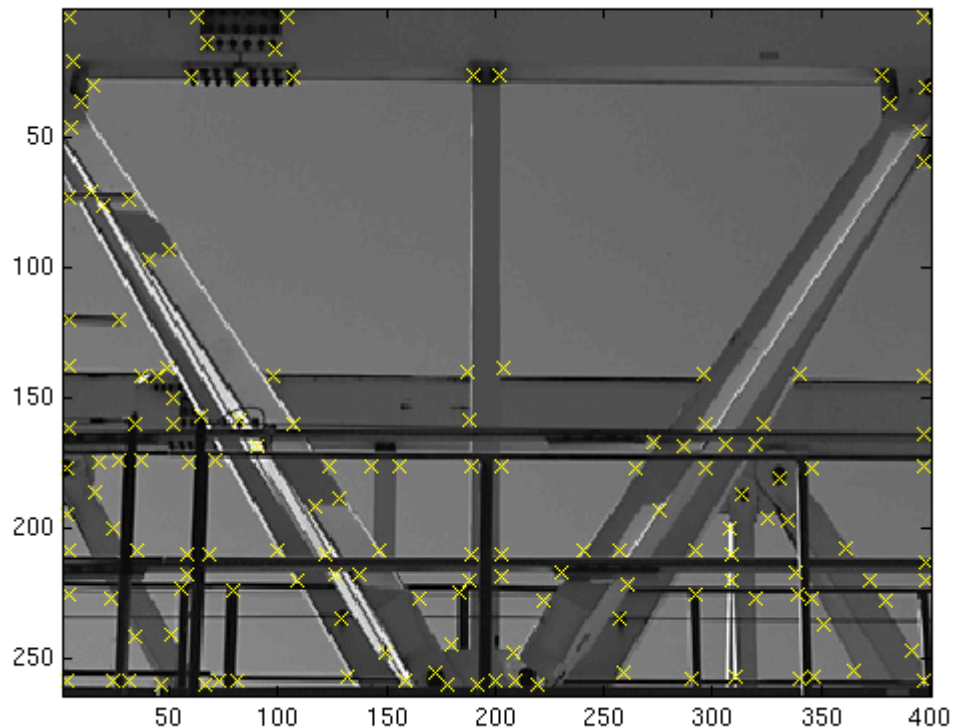
Feature detectors

Intensity based methods [Harris'88]



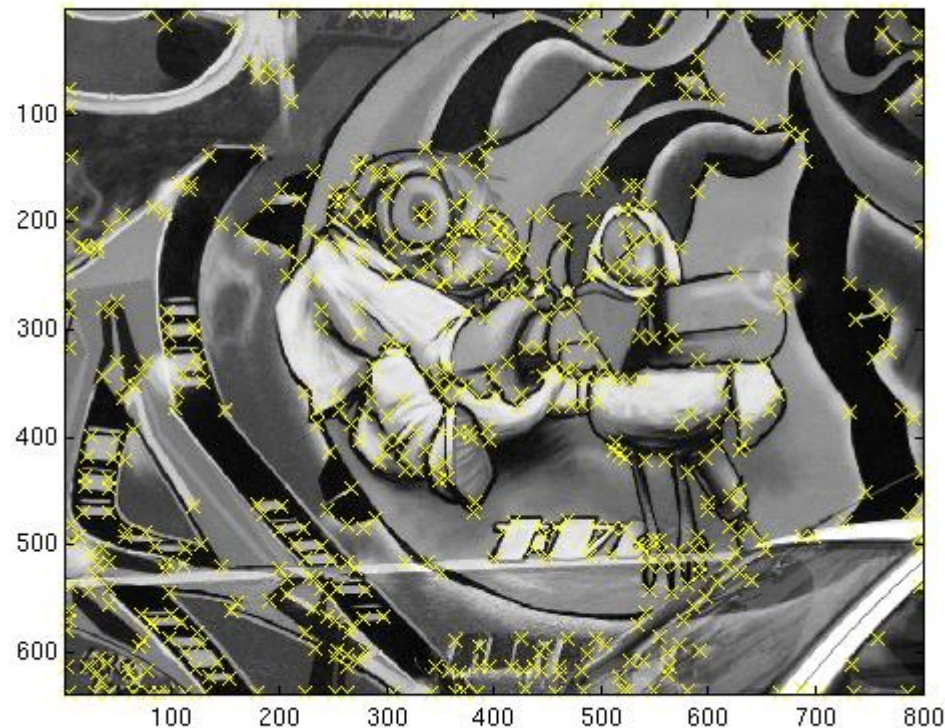
Feature detectors

Intensity based methods [Harris'88]



Feature detectors

Intensity based methods [Harris'88]

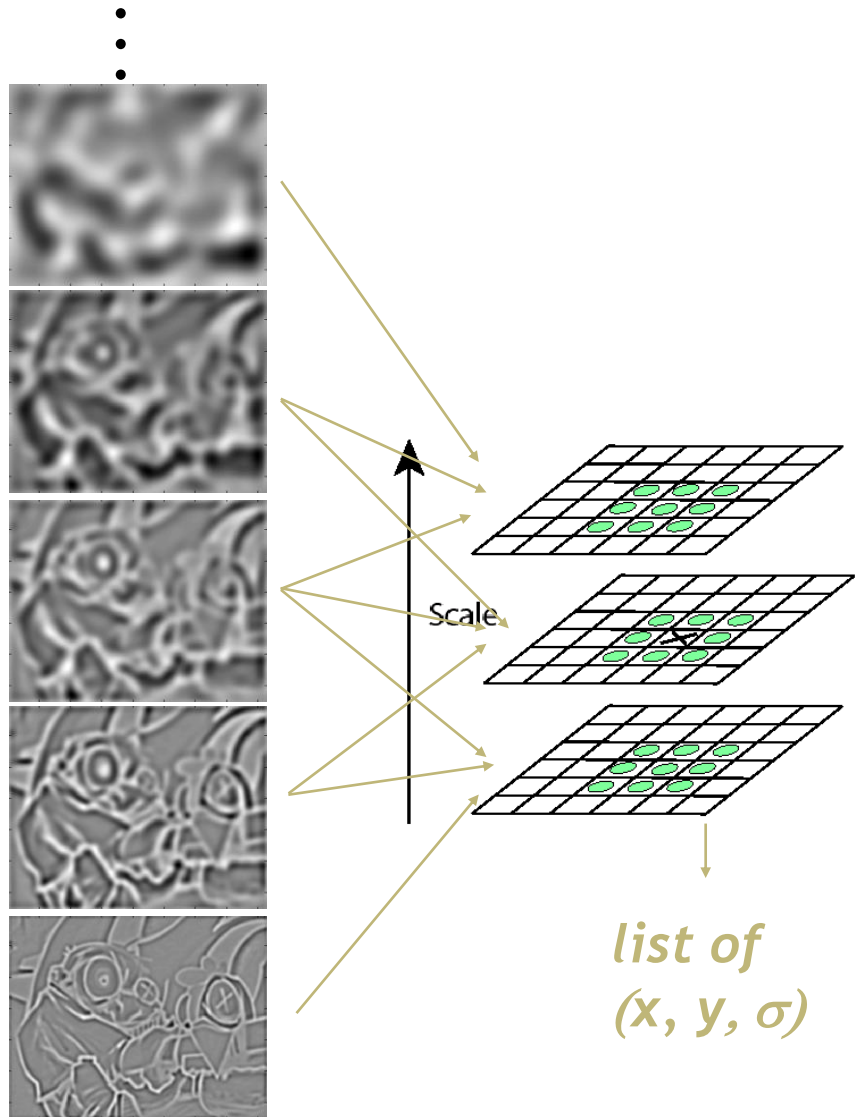
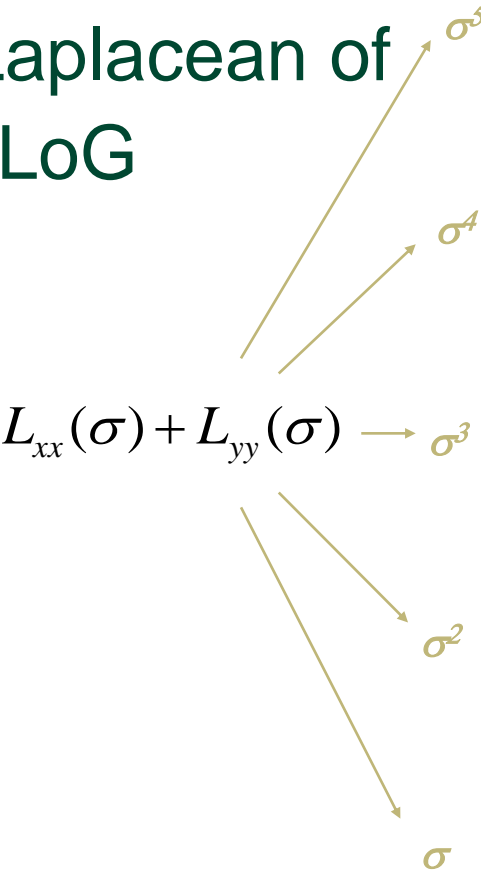


Scale-Space Methods

Scale invariant detectors

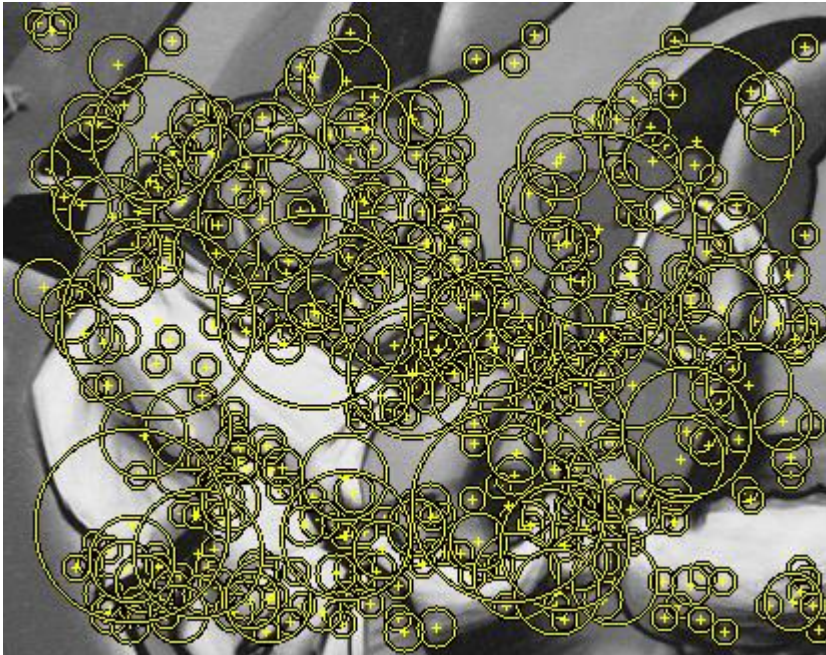
Laplacian of Gaussian

- Local maxima in scale space of Laplacean of Gaussian LoG



Scale invariant detectors

Laplacian of Gaussian



Scale Invariant detectors

Difference of Gaussian

- LoG \rightarrow diffusion equation \rightarrow derivative to scale

$$\frac{\partial L}{\partial s} = \vec{\nabla} \cdot \vec{\nabla} L = \Delta L = L_{xx}(\sigma) + L_{yy}(\sigma)$$

scale normalized
Laplacian

$$\sigma^2 (L_{xx}(\sigma) + L_{yy}(\sigma))$$

$$s = \sigma^2$$

$$\frac{\partial L}{\partial \sigma} \approx \frac{L(k\sigma) - L(\sigma)}{k\sigma - \sigma}$$

$$(k-1)\sigma^2 \Delta L \approx L(k\sigma) - L(\sigma)$$

$L(k\sigma)$



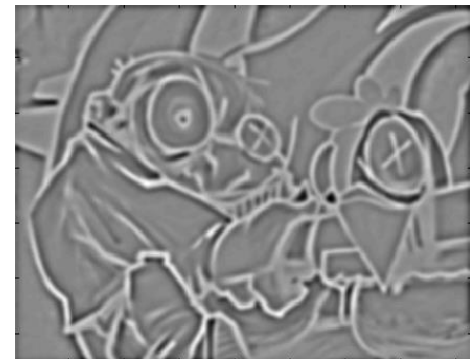
$L(\sigma)$



-

=

$L(k\sigma) - L(\sigma)$



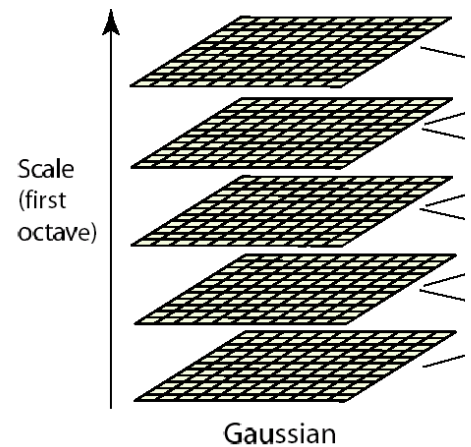
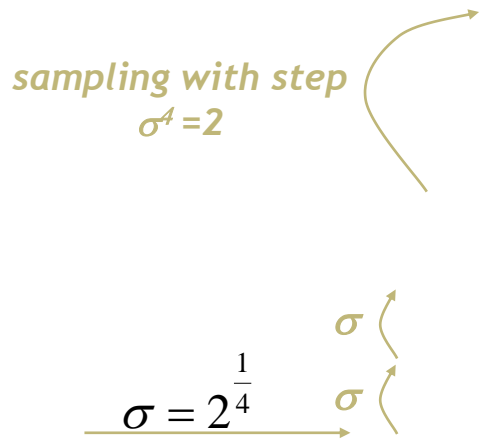
Scale invariant detectors

Difference of Gaussian

- Building scale space of Difference of Gaussian DoG
 - LoG \rightarrow diffusion equation \rightarrow derivative to scale



Original image



Scale invariant detectors

- Detecting multiscale points – thousands of interest points



$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha[\text{trace}(\mu(\sigma_I, \sigma_D))]$$

$$\sigma_I = 1.6 \cdot \sigma_D$$

Scale invariant detectors

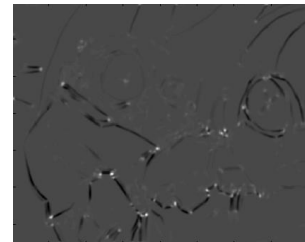
- Detecting multiscale points – thousands of interest points



$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))]$$

$$\sigma_I = 1.6 \cdot \sigma_D$$

 σ 

Computing Harris function

Scale invariant detectors

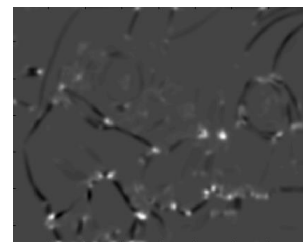
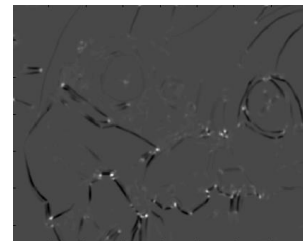
- Detecting multiscale points – thousands of interest points



$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))]^2$$

$$\sigma_I = 1.6 \cdot \sigma_D$$

 σ^2  σ 

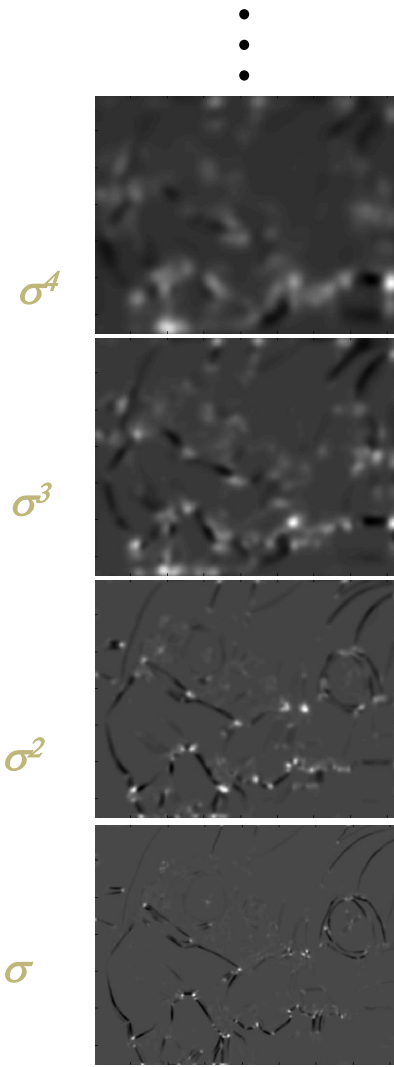
Computing Harris function

Scale invariant detectors

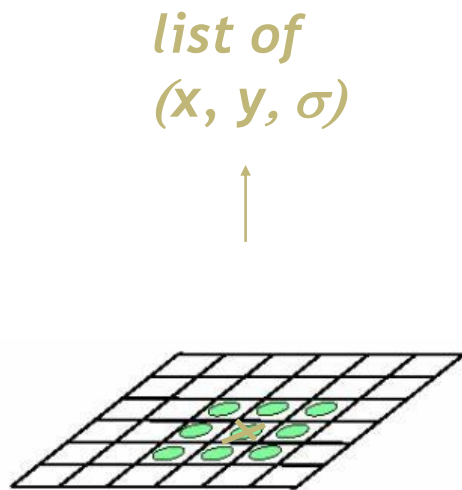
- Detecting multiscale points – thousands of interest points



$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$
$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))]^2$$
$$\sigma_I = 1.6 \cdot \sigma_D$$



Computing Harris function



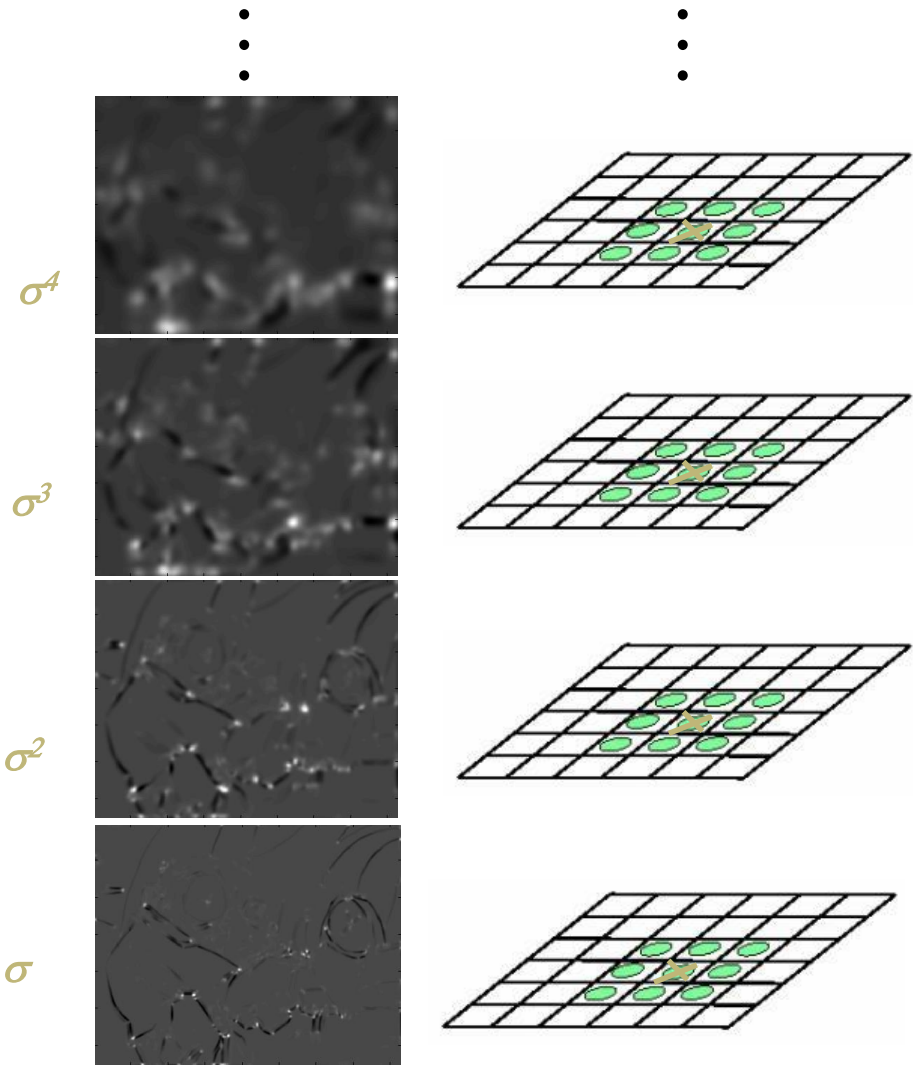
Detecting local maxima

Scale invariant detectors

- Detecting multiscale points – thousands of interest points



$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$
$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))]$$
$$\sigma_I = 1.6 \cdot \sigma_D$$



Computing Harris function

Detecting local maxima

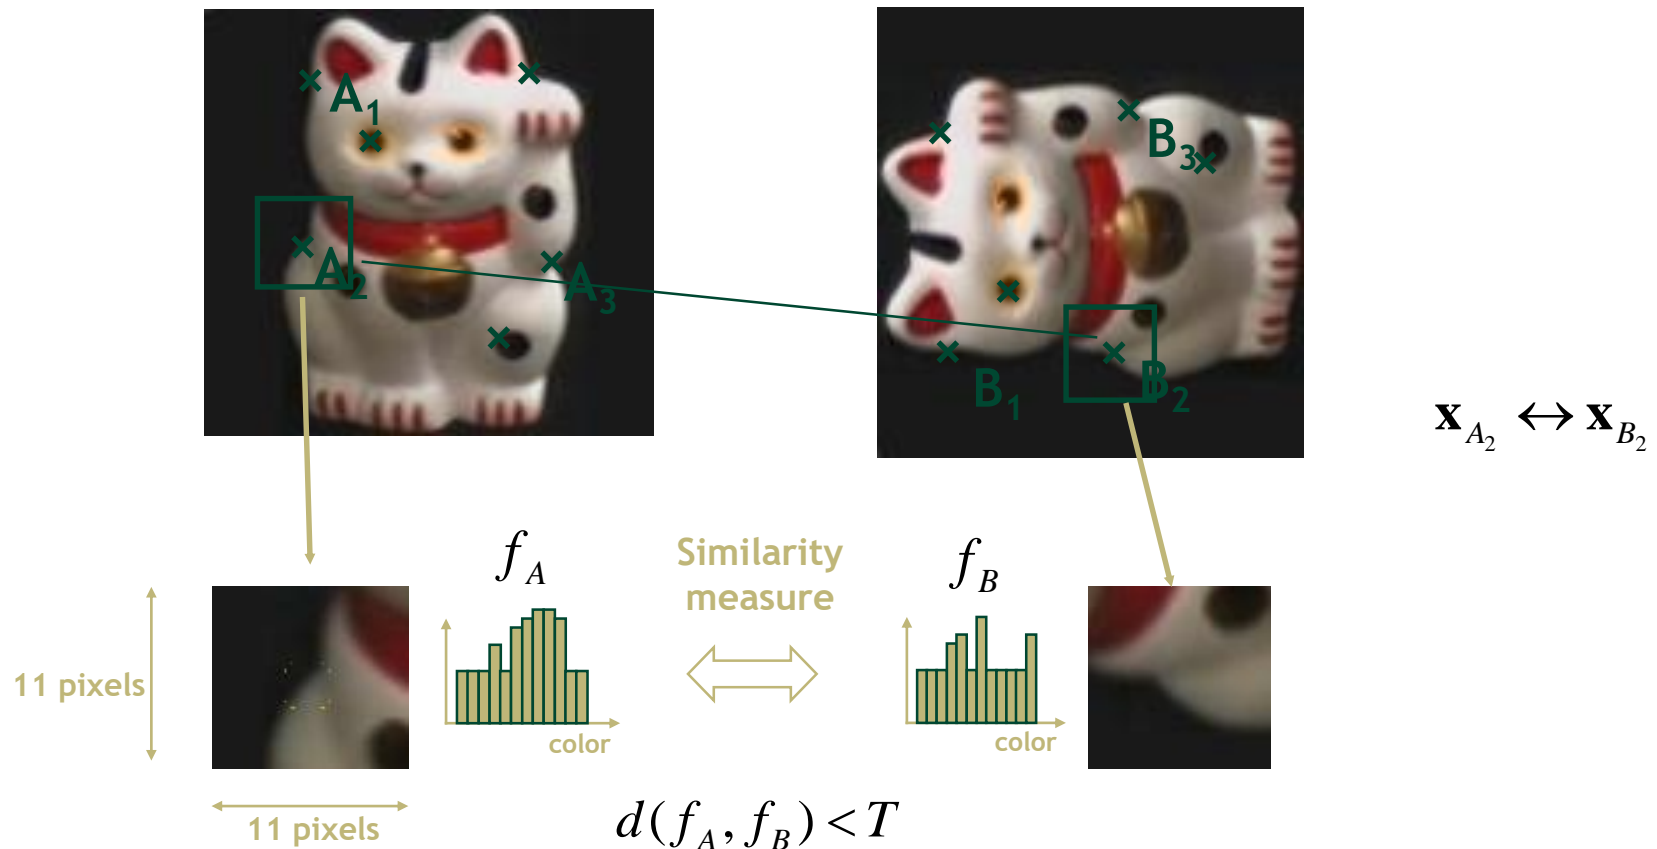
Summary

- Canny edge detection
- Hessian & second moment matrix, autocorrelation function
- Feature detection, types of features – methods to detect them.

Image transformation

Matching patches

- Extracting and matching patches



Descriptors history

Accuracy

- Normalized cross-correlation (NCC) [~ 1960s]
- Gaussian derivative-based descriptors
 - Differential invariants [Koenderink and van Doorn'87]
 - Steerable filters [Freeman and Adelson'91]
- Moment invariants [Van Gool et al.'96]
- SIFT [Lowe'99]
- Shape context [Belongie et al.'02]
- Gradient PCA [Ke and Sukthankar'04]

Efficiency

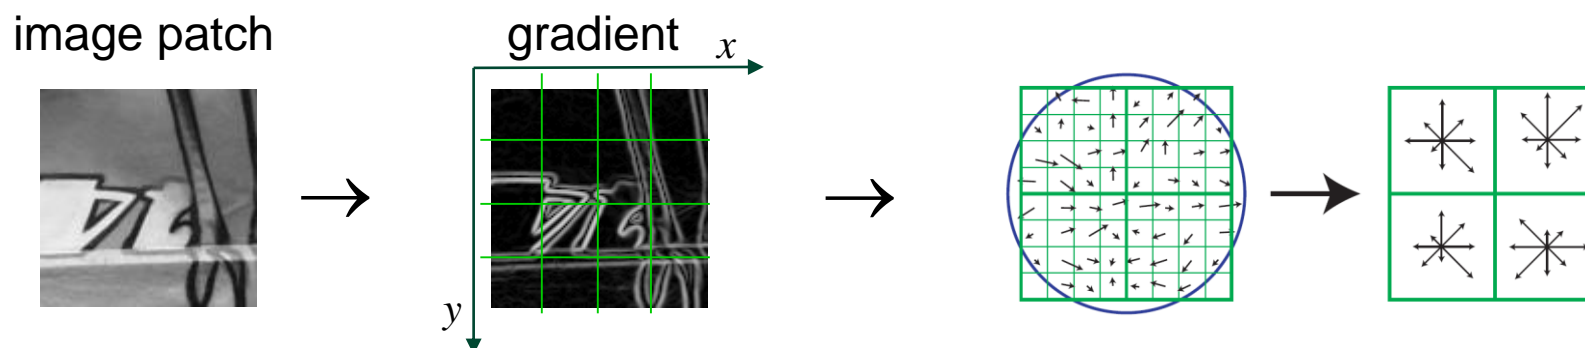
- SURF descriptor [Bay et al.'08]
- BRIEF [Calonder et al. 2010]

Machine learning

- Learning descriptors from image data [Brown et al 2010, ...]
- Neural Networks [Zagoruyko et al 2015, ...]

SIFT descriptor [Lowe'99]

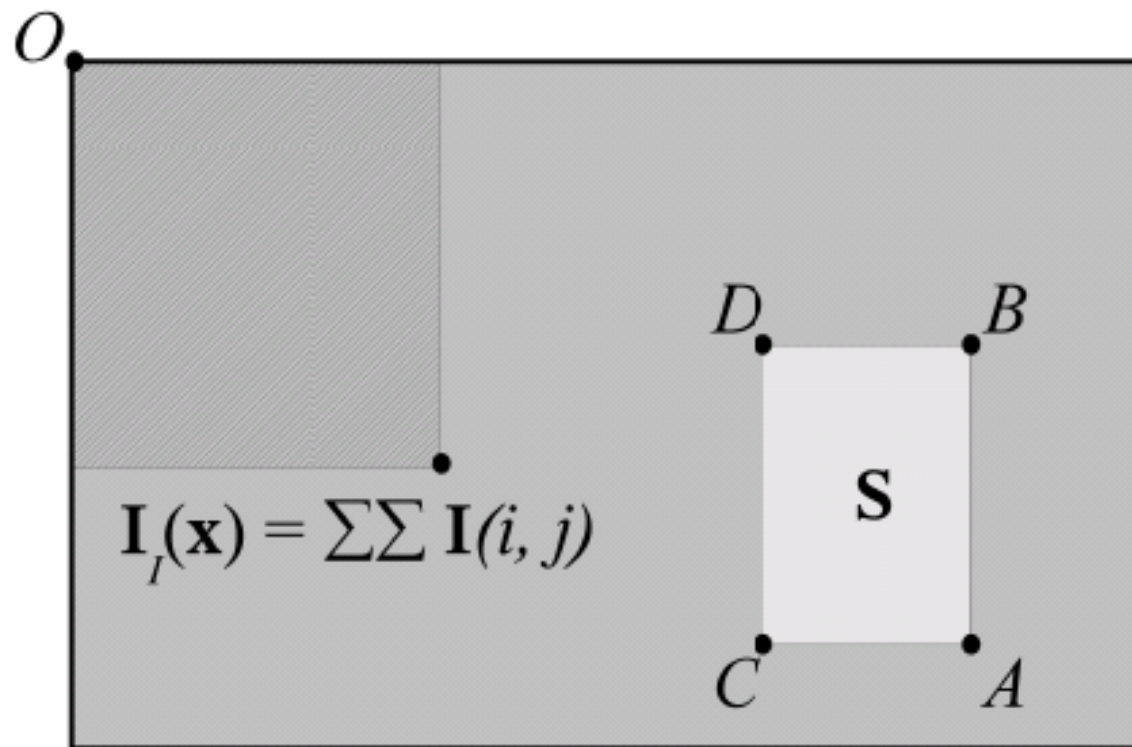
- Spatial binning and binning of the gradient orientation
- 4x4 spatial grid, 8 orientations of the gradient, dim 128
- Soft-assignment to spatial bins
- Normalization of the descriptor to norm one (robust to illumination)
- Comparison with Euclidean distance



SIFT Descriptor

- By far the most commonly used distinguished region descriptor:
 - fast
 - compact
 - works for a broad class of scenes
 - source code available
- Large number of ad hoc parameters) Enormous follow up literature on both “improvements” and improvements

The Integral image (Sum Table)



$$S = A - B - C + D$$

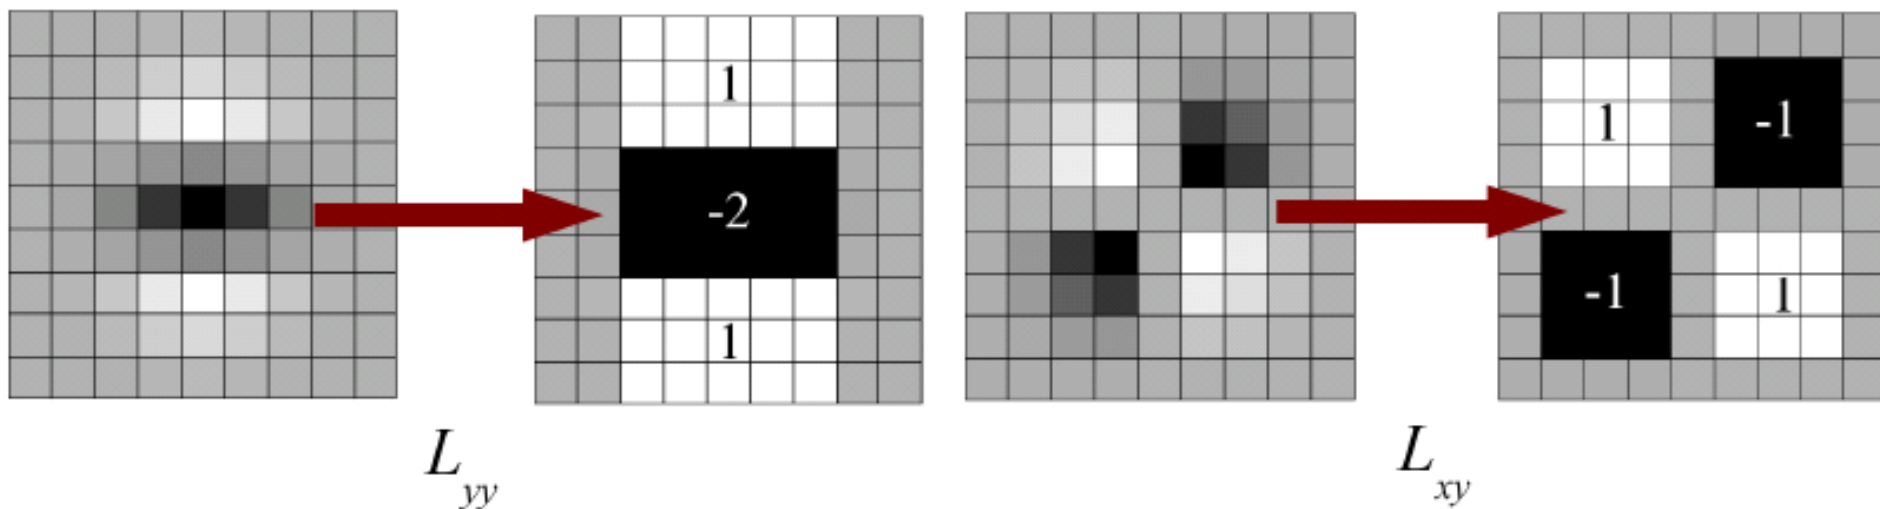
To calculate the sum in the DBCA rectangle, only 3 additions are needed

SURF Detection

- Hessian-based interest point localization:
- $L_{xx}(x,y,\sigma)$ is the convolution of the *Gaussian* second order derivative with the image

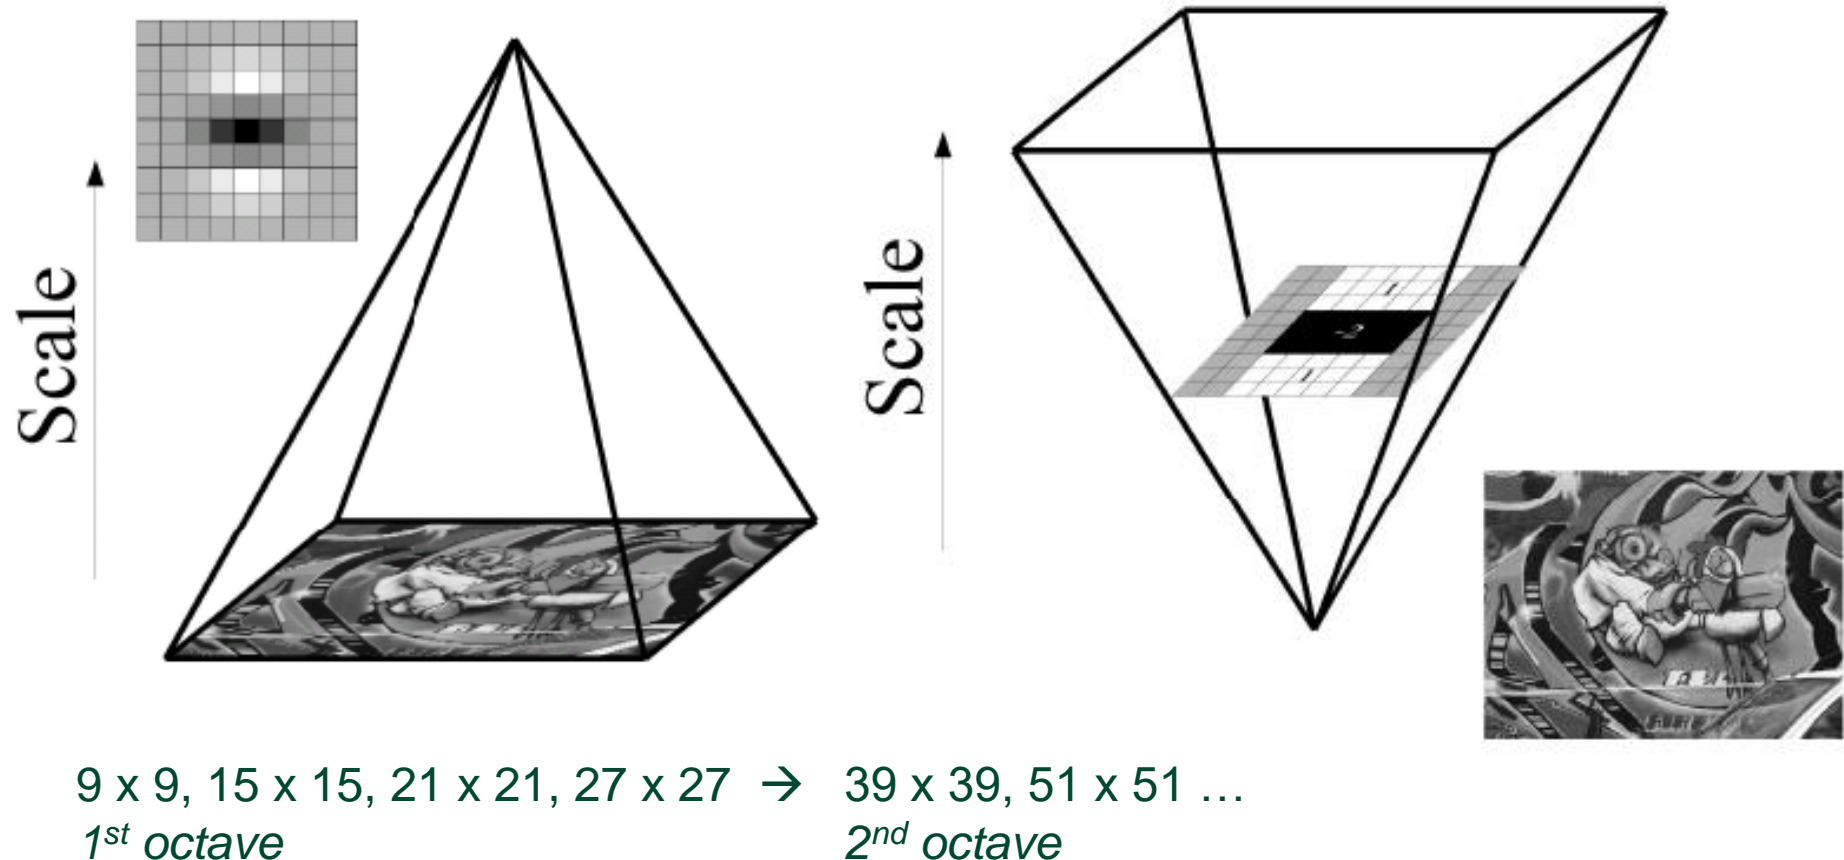
$$H = \begin{bmatrix} L_{xx} & L_{xy} \\ L_{xy} & L_{yy} \end{bmatrix}$$

- Approximate second order derivatives with box filters



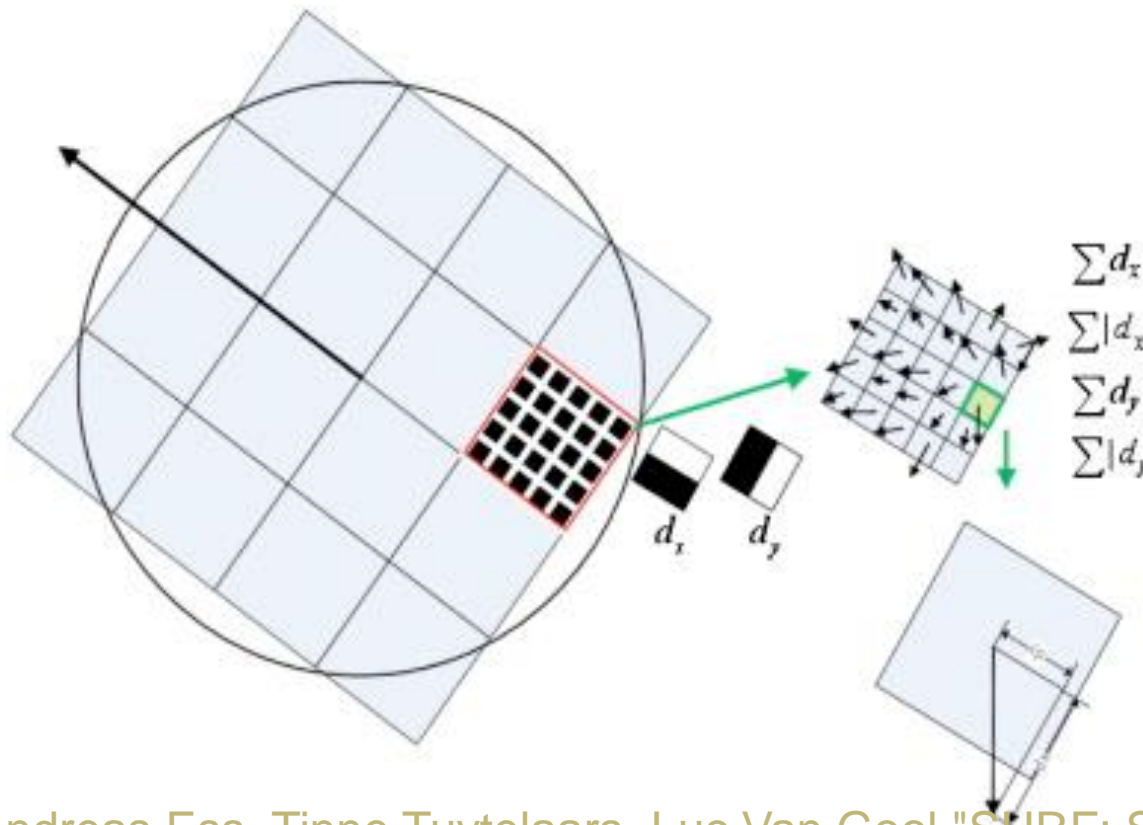
SURF Detection

- Scale analysis easily handled with the integral image



SURF: Speeded Up Robust Features

- Approximate derivatives with Haar wavelets
- Exploit integral images

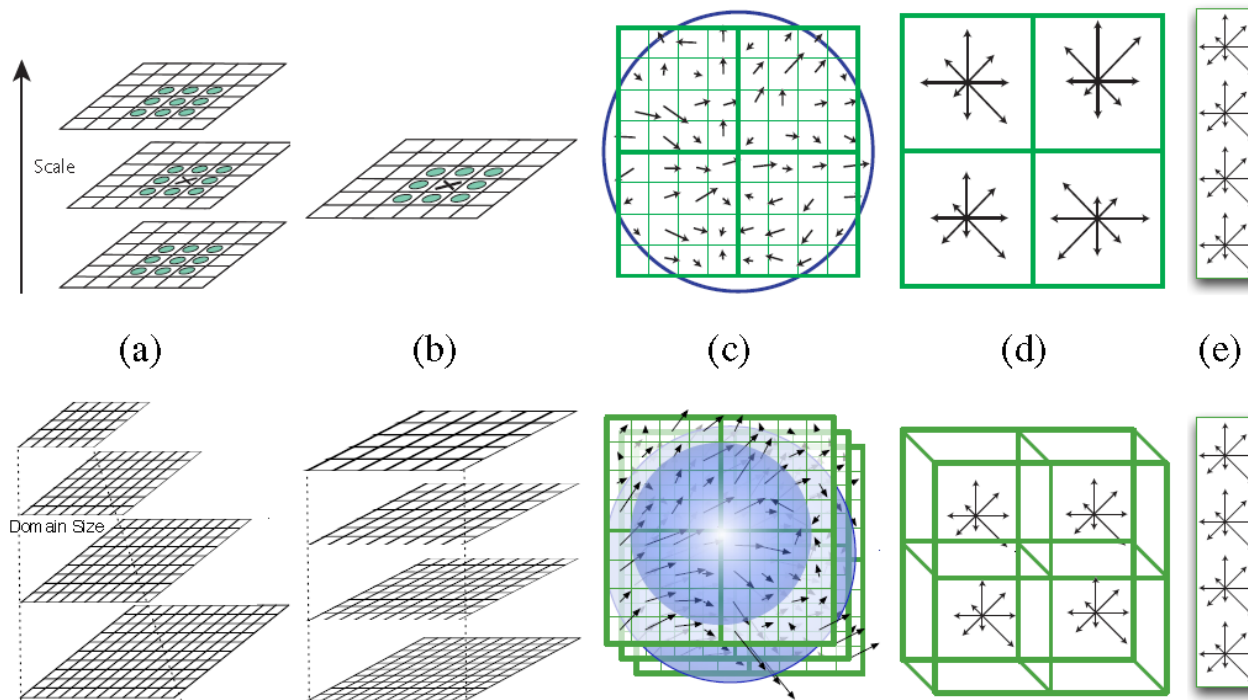


Herbert Bay, Andreas Ess, Tinne Tuytelaars, Luc Van Gool "SURF: Speeded Up Robust Features", Computer Vision and Image Understanding (CVIU), Vol. 110, No. 3, pp. 346--359, 2008

Questions: goo.gl/K61te5

DSP-SIFT

- Pooling SIFT descriptor across multiple scales



J. Dong, S. Soatto, Domain-Size Pooling in Local Descriptors: DSP-SIFT, CVPR2015
J. Dong, N Karianakis, D. Davis, J. Hernandez, J. Balzer, S. Soatto, Multi-View Feature Engineering and Learning, CVPR2015

Fast and compact descriptors

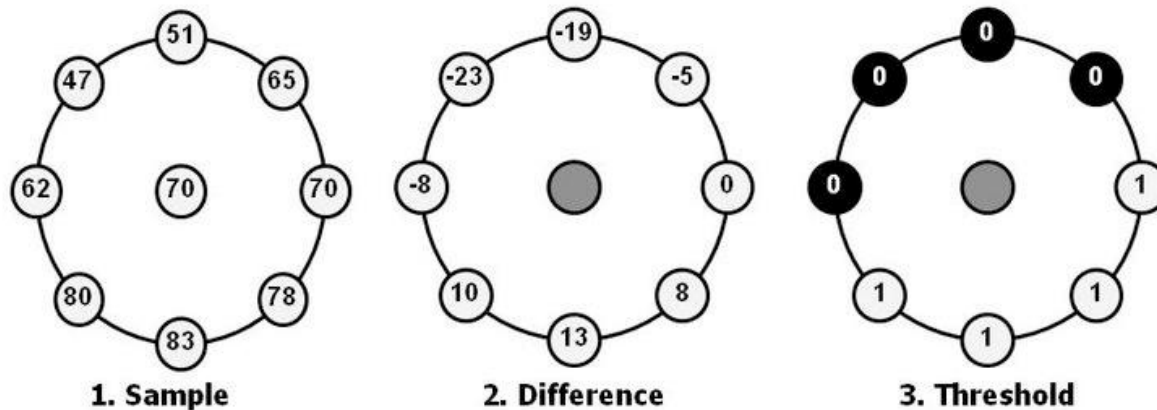
- Binary descriptors
- Comparison of pairs of intensity values
 - LBP
 - BRIEF, ORB, BRISK

LBP: Local Binary Patterns

- First proposed for texture recognition in 1994.

The value of the LBP code of a pixel (x_c, y_c) is given by:

$$LBP_{P,R} = \sum_{p=0}^{P-1} s(g_p - g_c) 2^p \quad s(x) = \begin{cases} 1, & \text{if } x \geq 0; \\ 0, & \text{otherwise.} \end{cases}$$



$$1 \cdot 1 + 1 \cdot 2 + 1 \cdot 4 + 1 \cdot 8 + 0 \cdot 16 + 0 \cdot 32 + 0 \cdot 64 + 0 \cdot 128 = 15$$

4. Multiply by powers of two and sum

T. Ojala, M. Pietikäinen, and D. Harwood (1994), "Performance evaluation of texture measures with classification based on Kullback discrimination of distributions", ICPR 1994, pp.582-585.

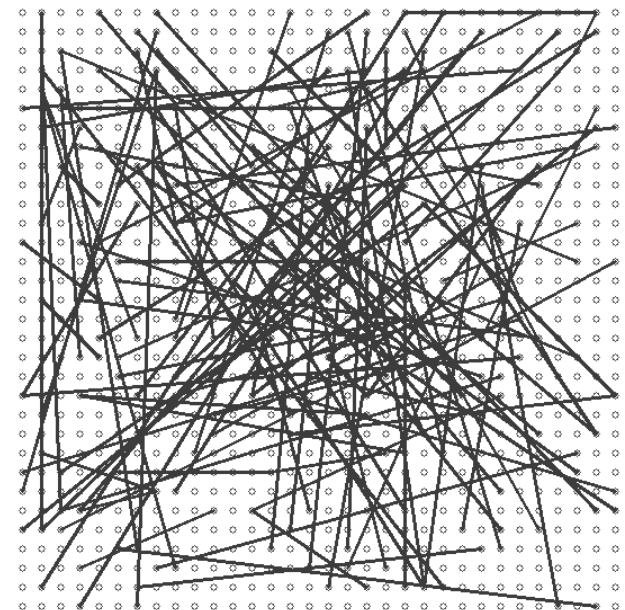
M Heikkilä, M Pietikäinen, C Schmid, Description of interest regions with LBP, Pattern recognition

42 (3), 425-436
Questions: goo.gl/K6ite5

BRIEF:

Binary Robust Independent Elementary Features

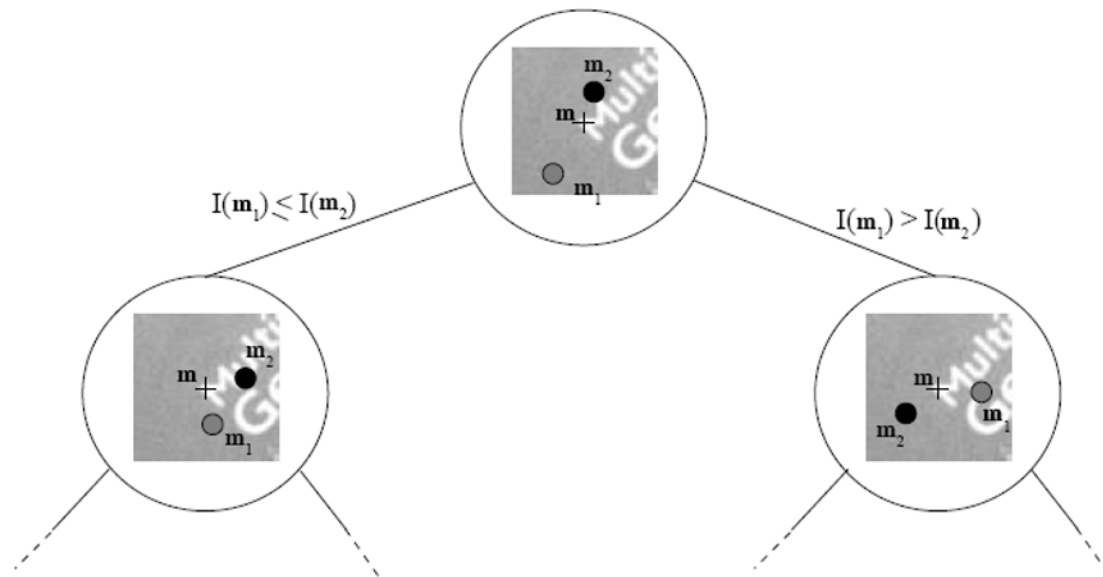
- Random selection of pairs of intensity values.
- Fixed sampling pattern of 128, 256 or 512 pairs.
- Hamming distance to compare descriptors (XOR).



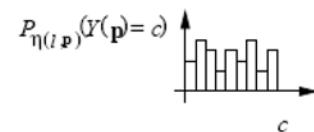
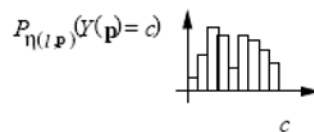
M. Calonder, V. Lepetit, C. Strecha, P. Fua, BRIEF: Binary Robust Independent Elementary Features, 11th European Conference on Computer Vision, 2010.

Randomized Decision Tree

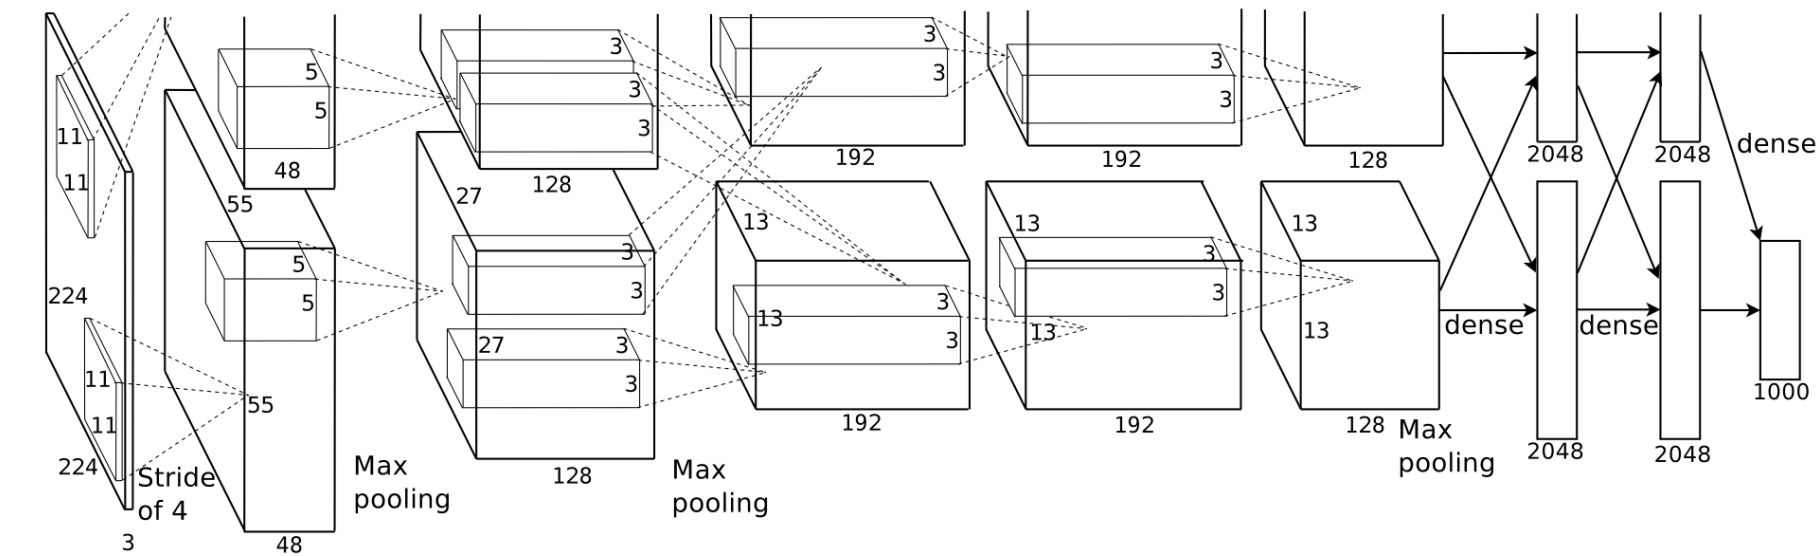
- Compare intensity of pairs of pixels
- In construction, pick pairs randomly
 - Insert all training examples into tree
 - Distribution at leaves is descriptor for the particular feature



Lepetit, Laguerre and Fua.
Randomized Trees for Real-Time
Keypoint Matching, CVPR 2005



Convolutional Neural Networks

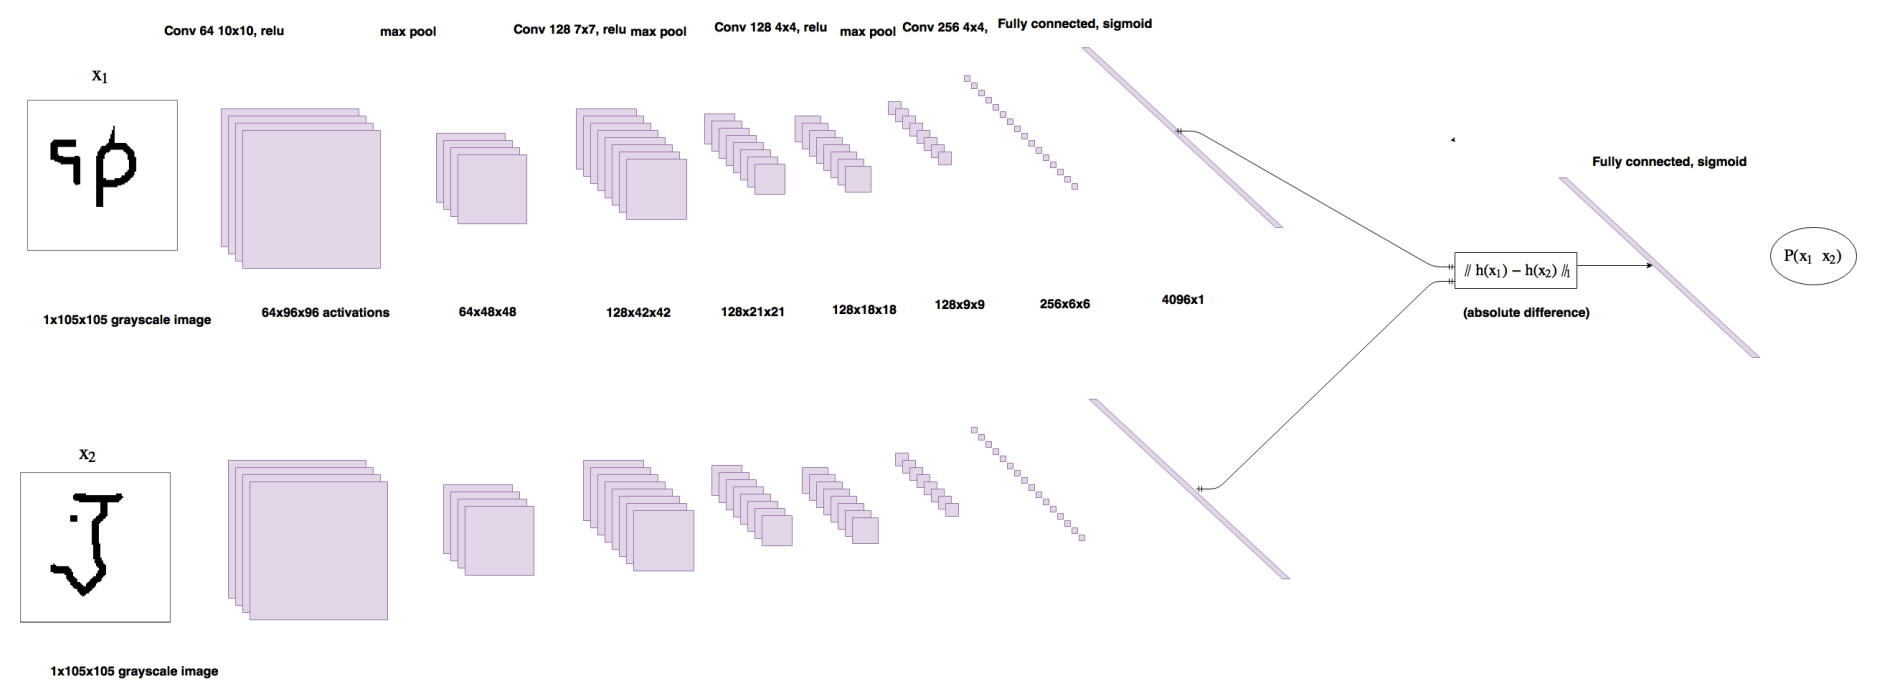


Model	Top-1	Top-5
<i>Sparse coding [2]</i>	47.1%	28.2%
<i>SIFT + FVs [24]</i>	45.7%	25.7%
CNN	37.5%	17.0%

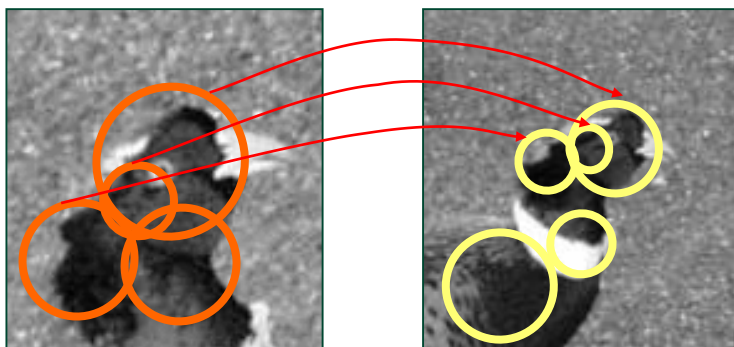
Alex Krizhevsky, Ilya Sutskever Geoffrey E. Hinton, ImageNet Classification with Deep Convolutional Neural Networks , NIPS 2012

Convolutional neural networks

- Siamese networks

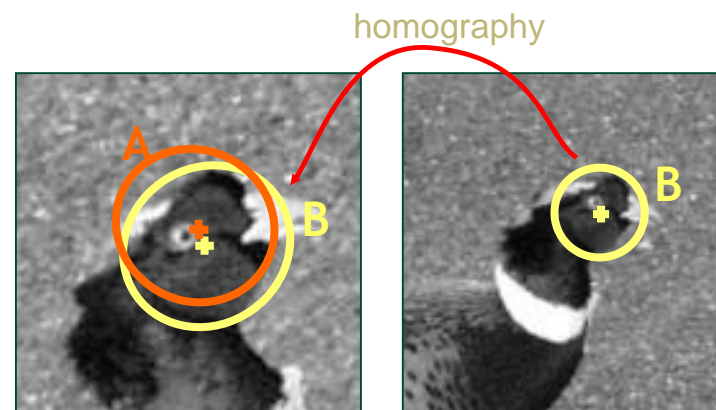


Detector evaluations



$$\text{precision} = \frac{\# \text{correct matches}}{\# \text{all matches}}$$

$$\text{recall} = \frac{\# \text{correct matches}}{\# \text{ground truth correspondences}}$$

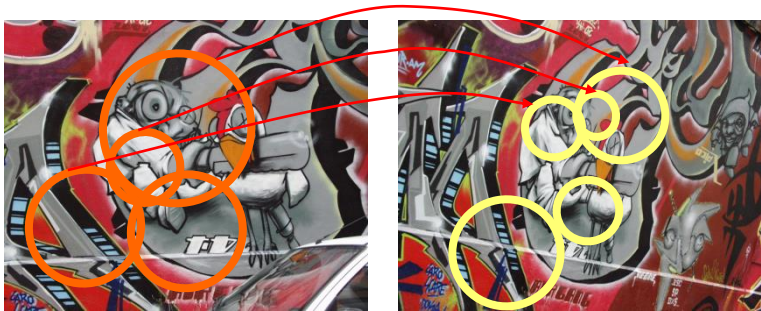


Two points are correctly matched if
 $T=40\%$

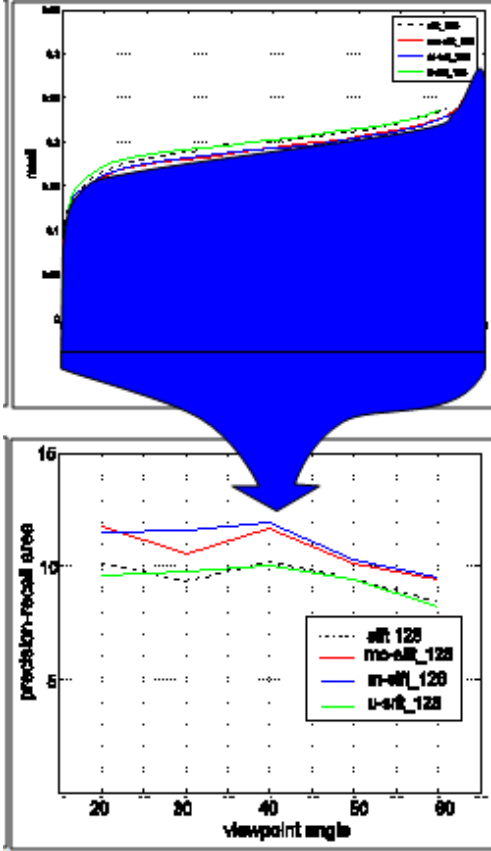
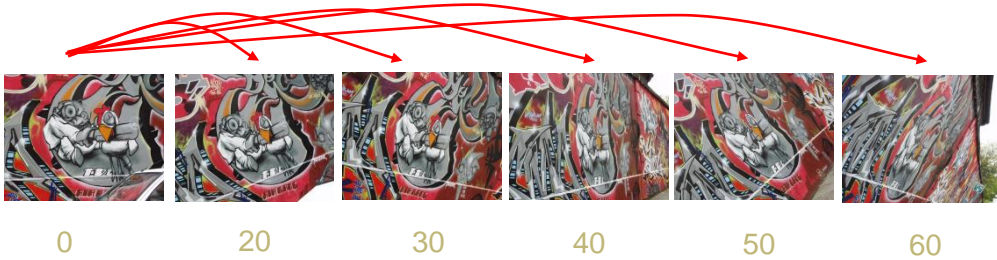
$$\frac{A \cap B}{A \cup B} > T$$

Matching test

Precision-recall area



$$\text{precision} = \frac{\# \text{correct matches}}{\# \text{all matches}}$$
$$\text{recall} = \frac{\# \text{correct matches}}{\# \text{ground truth correspondences}}$$



Summary

- Feature detection
 - Edge
 - Interest points
- Feature description
 - SIFT
 - SURF
 - Binary
 - Neural Networks
 - Etc.
- Evaluations
 - Measures and protocol