



Motion

Motion Estimation

- Given a sequence of images we might ask
 - What are the moving objects in the scene?
 - What sort of motion are they undergoing?
 - Where will they be in the future?
- To answer these questions we need to measure the motion
- There are many problems in motion estimation
- Often the motion is ambiguous
- Image sequences contain a lot of data - efficiency is a concern
- Many interesting tasks involve complex motion - e.g. facial expression analysis

 $I(x,y,t)$  $I(x,y,t+1)$

Motion

- Motion and stereo are closely related
 - “Correspondence problem”, or “visual correspondence”: what went where?
 - Can be solved sparsely or densely
 - Except for wide-baseline stereo, in both cases sparse solutions are now rare
- Comparison:
 - Stereo uses a 1D label set, motion has 2D
 - Stereo involves bigger changes in appearance
 - If motion is small, tracking is an easy way
 - Motion has an elegant formulation as a continuous problem

Challenges

- Figure out which features can be tracked
- Efficiently track across frames
- Some points may change appearance over time (e.g., due to rotation, moving into shadows, etc.)
- Drift: small errors can accumulate as appearance model is updated
- Points may appear or disappear: need to be able to add/delete tracked points

Simple Techniques

Motion Difference

- Take two images from a sequence
- Compute the change in brightness at each pixel in the image
- Threshold


Background Models

- Find the average brightness at each pixel over a sequence
- Use the difference between the current frame and the average to find moving objects



Simple Techniques

- Area-based matching can also be used
- We take a template from the first image
- This is then compared to points in the second image to find corresponding regions
- This uses a 'distance measure' to compare patches



Motion Field and Optical Flow

The Motion Field


- “assigns a velocity vector to each point in the image”
- Tells us how the position of the *image* of the corresponding *scene point* changes over time
- Can be computed from the *scene* to tell us about the *image*

Optical Flow

- The “apparent motion of the brightness pattern” in an image
- Ideally it will be the same as the motion field, but this is not always the case
- Can be computed from the *image*, to tell us about the *scene*


Optical Flow \neq Motion Field

A Moving light



- The *image* changes so there is optical flow
- The *scene* objects do not move so there is no motion field

A Rotating Sphere





- The *scene* object moves, so there is motion field
- The *image* does not change, so there is no optic flow

Optical flow



- Our “data cost” will be the difference between the old pixel intensity and the new pixel intensity
 - New and old pixel are related by the motion
- Instead of search, we can directly solve for a data cost
 - We consider a series of images as samples of a function $I(x,y,t)$
 - What are we assuming?
- No explicit search over (many) labels in contrast to stereo

Optical Flow is Ambiguous

- Consider the two images below:



- Two possible fields (of



So, optical flow

- Is not always what we want to compute
- Cannot be determined without ambiguity
- But it is all that we can compute from the images
- This means we need to make assumptions to find a *reasonable* flow field estimate

Brightness Constancy

Brightness constancy assumption:



$$I(x, y, t) = I(x + u, y + v, t + 1)$$

- $I(x, y, t)$ is the brightness of the image at location (x, y) and time t
- (u, v) is the motion field at location (x, y) and time t
- This assumption is true apart from the effects of lighting (including shadows, reflections, and highlights)

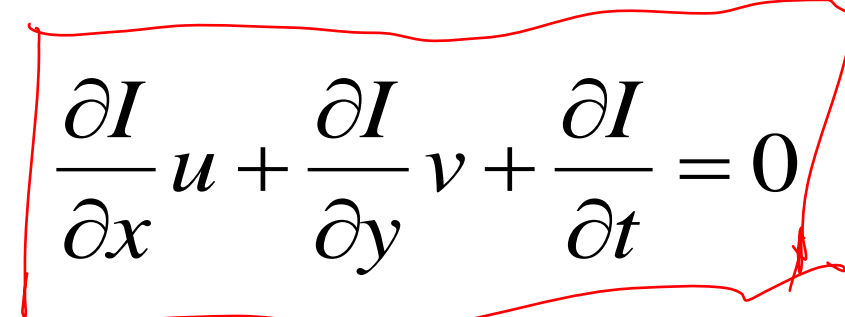
Brightness Constancy

Another way to express brightness constancy is that

$$\frac{dI(x,y,t)}{dt} = 0$$

- This says that the image doesn't change over time - it just moves about

$$\frac{\partial I}{\partial x} \frac{dx}{dt} + \frac{\partial I}{\partial y} \frac{dy}{dt} + \frac{\partial I}{\partial t} = 0$$


$$\frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \frac{\partial I}{\partial t} = 0$$

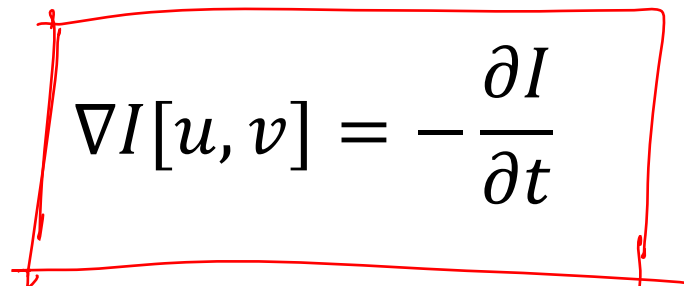
Brightness Constancy

$$\frac{\partial I(x, y, t)}{\partial x} u + \frac{\partial I(x, y, t)}{\partial y} v + \frac{\partial I(x, y, t)}{\partial t} = 0$$

Image derivative
in x direction

Image derivative
in y direction

Image derivative
in t direction



$$\nabla I[u, v] = -\frac{\partial I}{\partial t}$$

The Aperture Problem

There is no solution to the equation

We only measure the projection of the true motion (u, v) on the intensity gradient, which makes it ambiguous

- We can determine the component of flow in the same direction as the image intensity gradient
- We cannot determine the component of flow perpendicular to it
- This is the *aperture problem*



Flow Smoothness

We need another constraint to find a unique solution

- This is the constraint that the flow field is smooth
- Neighbouring pixels in the image should have similar optical flow

We want u and v to have low variation

- We can do this by trying to set

$$(u - \bar{u}) = 0, \quad (v - \bar{v}) = 0$$

- So u and v are equal to the average of their neighbouring values

Smoothness assumption







- One-dimensional example –linearization
 - Estimate displacement u using derivative
 - Two functions $f(x)$ and $g(x)=f(x-u)$
 - Taylor series expansion

$$f(x-u) = f(x) - u f'(x) + E$$

f' denotes derivative
 - write difference as

$$f(x)-g(x) = u f'(x) + E$$
 - Discarding higher order terms

$$\delta = (f(x)-g(x))/f'(x)$$
 - works only for small u
- We need more than 1 pixel
- Each pixel defines linear constraint on possible (u,v) displacement
 - For set of pixels with same displacement combine constraints to get estimate
 - For pixels with different displacements, somehow identify that case



Squared Errors

We now have three error terms

- If we square them then the error is always positive, and we can look for a minimum
- A weighting term, λ , balances the influence of the brightness and smoothness errors

The squared error term is

- To minimise: take derivatives with respect to u and v , set to 0, then solve

$$\lambda \left(\frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \frac{\partial I}{\partial t} \right)^2 + (u - \bar{u})^2 + (v - \bar{v})^2$$

Minimisation

$$e = \lambda \left(\frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \frac{\partial I}{\partial t} \right)^2 + (u - \bar{u})^2 + (v - \bar{v})^2$$

$$\frac{\partial e}{\partial u} = 2\lambda \frac{\partial I}{\partial x} \left(\frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \frac{\partial I}{\partial t} \right) + 2(u - \bar{u}) = 0$$

$$\frac{\partial e}{\partial v} = 2\lambda \frac{\partial I}{\partial y} \left(\frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \frac{\partial I}{\partial t} \right) + 2(v - \bar{v}) = 0$$

Solving the two equations
gives

$$u = \bar{u} - \lambda \frac{\frac{\partial I}{\partial x} \bar{u} + \frac{\partial I}{\partial y} \bar{v} + \frac{\partial I}{\partial t}}{1 + \lambda \left(\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2 \right)}$$


$$v = \bar{v} - \lambda \frac{\frac{\partial I}{\partial x} \bar{u} + \frac{\partial I}{\partial y} \bar{v} + \frac{\partial I}{\partial t}}{1 + \lambda \left(\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2 \right)}$$

But we need to know $\sim u$ and
 $\sim v$ to compute u and v

Iterative solution:

- Estimate u and v
- Then compute the averages,
 $\sim u$ and $\sim v$
- Then make a new estimate
of u and v
- Then make a new estimate
of $\sim u$ and $\sim v$
- etc...

Computing the Optical Flow



Gradients:

$$dI/dx = (b+d+f+h) - (a+c+e+g)$$

$$dI/dy = (a+b+e+f) - (c+d+g+h)$$

$$dI/dt = (e+f+g+h) - (a+b+c+d)$$

$$u = \bar{u} - \lambda \frac{\partial I}{\partial x} \frac{\bar{u} \frac{\partial I}{\partial x} + \bar{v} \frac{\partial I}{\partial y} + \frac{\partial I}{\partial t}}{1 + \lambda \left(\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2 \right)}$$

$$v = \bar{v} - \lambda \frac{\partial I}{\partial y} \frac{\bar{u} \frac{\partial I}{\partial x} + \bar{v} \frac{\partial I}{\partial y} + \frac{\partial I}{\partial t}}{1 + \lambda \left(\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2 \right)}$$

Computing the Optical Flow

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	2	1	0	0	0	0
0	0	(x,y,t)	5	1	0	0	0	0
0	0	2	6	4	0	0	0	0
0	0	0	3	1	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

5	5
6	4

$$\cdot \frac{\frac{\partial I}{\partial y} + \frac{\partial I}{\partial t}}{\sqrt{1 + \left(\frac{\partial I}{\partial y}\right)^2}}$$

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	1	2	1	0	0	0
0	0	1	4	5	4	1	0	0
0	0	0	4	6	3	1	0	0
0	0	0	1	2	2	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

4	5
4	6

$$\cdot \frac{\frac{\partial I}{\partial y} + \frac{\partial I}{\partial t}}{\sqrt{1 + \left(\frac{\partial I}{\partial y}\right)^2}}$$

$$dI/dx = (5+4+5+6)-(5+6+4+4)$$

$$dI/dy = (5+5+4+5)-(6+4+4+6)$$

$$dI/dt = (5+5+6+4)-(4+5+4+6)$$

Lukas-Kanade

- If there is a single translational motion (u,v)
 - In a window, or over the entire image
- At each pixel, the OFCE (optical flow constraint equation) says:

- the spatial and temporal gradient $\nabla I[u, v] = -\frac{\partial I}{\partial t}$
- These are the observations $I_x(x_i, y_i) \cdot u + I_y(x_i, y_i) \cdot v = -I_t(x_i, y_i)$

$$\begin{bmatrix} I_x(x_1, y_1) & I_y(x_1, y_1) \\ I_x(x_2, y_2) & I_y(x_2, y_2) \\ \vdots & \vdots \\ I_x(x_n, y_n) & I_y(x_n, y_n) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -I_t(x_1, y_1) \\ -I_t(x_2, y_2) \\ \vdots \\ -I_t(x_n, y_n) \end{bmatrix} \quad S \cdot \begin{bmatrix} u \\ v \end{bmatrix} = -T$$

- We can use least squares to solve this

Lukas-Kanade

- Least square solution to

$$\begin{bmatrix} I_x(x_1, y_1) & I_y(x_1, y_1) \\ I_x(x_2, y_2) & I_y(x_2, y_2) \\ \vdots & \vdots \\ I_x(x_n, y_n) & I_y(x_n, y_n) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -I_t(x_1, y_1) \\ -I_t(x_2, y_2) \\ \vdots \\ -I_t(x_n, y_n) \end{bmatrix} \quad S \cdot \begin{bmatrix} u \\ v \end{bmatrix} = -T$$

$$S^t S \cdot \begin{bmatrix} u \\ v \end{bmatrix} = -S^t T$$

second
order
matrix

$$S^t S = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} \quad S^t T = \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$$Ad = b \xRightarrow{LS} \min_d \|Ad - b\|^2$$

Lukas-Kanade

$$S^t S \cdot \begin{bmatrix} u \\ v \end{bmatrix} = -S^t T \qquad \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

- When is this solvable? i.e., what are good points to track?
 - $S^T S$ should be invertible
 - $S^T S$ should not be too small due to noise
 - eigenvalues e_1, e_2 of $S^T S$ should not be too small
 - $S^T S$ should be well-conditioned
 - e_1 / e_2 should not be too large (e_1 = larger eigenvalue)
- Inverting 2x2 matrix has a closed form solution
- Unless spatial gradients are parallel
 - Multiple copies of the same gradient give singular matrix
 - We need textured regions
- Best if gradients are orthogonal e.g. corners

Computing the Optical Flow


The algorithm is iterative

- We start with an initial estimate
- We refine it over a series of cycles
- We need an initial estimate
- We also need to know when to stop


Initialisation

- We can start with an estimate of u and v of 0 everywhere
 - From coarse to fine
- Stop when the results at iteration n and $n+1$ are very similar
 - This is when the algorithm converges
 - Can we be sure it will?


Computing the Optical Flow




1 Iteration




2 Iterations




3 Iterations



5 Iterations



10 Iterations



15 Iterations

Template tracking

$$E(u, v) = \sum [I(x + u, y + v) - T(x, y)]^2$$



current frame



template
(model)

u, v = hypothesized location of
template in current frame

Template tracking

$$\begin{aligned} E(u, v) &= \sum [I(x+u, y+v) - T(x, y)]^2 \\ &\approx \sum [I(x, y) + uI_x(x, y) + vI_y(x, y) - T(x, y)]^2 \quad \text{First order approx} \\ &= \sum [uI_x(x, y) + vI_y(x, y) + D(x, y)]^2 \end{aligned}$$

Take partial derivs and set to zero

$$\frac{\delta E}{du} = \sum [uI_x(x, y) + vI_y(x, y) + D(x, y)] I_x(x, y) = 0$$

$$\frac{\delta E}{dv} = \sum [uI_x(x, y) + vI_y(x, y) + D(x, y)] I_y(x, y) = 0$$

Form matrix equation

$$\sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \sum \begin{bmatrix} I_x D \\ I_y D \end{bmatrix} \quad \Rightarrow \quad \boxed{\text{solve via least-squares}}$$

Template tracking

- Assumption of constant flow (pure translation) for all pixels in a larger window is unreasonable for long periods of time



- we can generalize Lucas-Kanade approach to other 2D parametric motion models (like affine or projective) by introducing a “warp” function W

$$E(u, v) = \sum [I(x+u, y+v) - T(x, y)]^2 \xrightarrow{\text{generalize}} \sum [I(W([x, y]; P)) - T([x, y])]^2$$

Affine motion

- An affine motion maps between two arbitrary triangles

$$u(x, y) = a_1 + a_2x + a_3y$$

$$v(x, y) = a_4 + a_5x + a_6y$$

- Generalisation of translational motion where $a_2 = a_3 = a_5 = a_6 = 0$

- Using OFCE we get:

$$I_x(x_i, y_i) \cdot (a_1 + a_2x + a_3y) + \\ I_y(x_i, y_i) \cdot (a_4 + a_5x + a_6y) = -I_t(x_i, y_i)$$

- More complex optimization

- 6 parameters to find instead of 2
- Still can use least squares

Motion estimates

- From motion estimates you can compute the focus of expansion
 - Direction of heading, relative to camera by intersecting motion rays
 - clustering of moving pixels e.g. Hough transform voting scheme
- Frequently use for registering images
- Object tracking
- Useful for robotic applications i.e. efficient