

Pattern Recognition

Krystian Mikolajczyk

Blackboard

Model Fitting

Data Representation

- Non-parametric representation
 - Collection of samples, data points

 x_1, y_1 x_2, y_2


.

.


.

.

.

 x_N, y_N 


Data Representation



Data Representation

- Parametric representation of a model

$$y=ax+b$$



$$(x-a)^2+(y-b)^2=r^2$$

RANSAC

RANSAC (Fischler and Bolles 1981)


- The RANSAC is an algorithm for robust fitting of models in the presence of many data outliers.
 - Given a fitting problem with parameters, estimate the model parameters from the data.
- Assumptions:
 - The model parameters can be estimated from m data items out of total $M \gg m$ items
 - Fit/non-fit can be established for every data item
 - The probability of a randomly selected data item being part of a good model is p_g
 - The probability that the algorithm will exit without finding a good fit (if one exists) is p_{fail}

Example: find parameters of a line going through the points

$$ax + by + c = 0 \quad y = ux + v$$

Need 2 points to uniquely define a line ($m=2$)

$$P_1(x_1, y_1), P_2(x_2, y_2) \quad u = \frac{y_1 - y_2}{x_1 - x_2} \quad v = \frac{x_2 y_1 - x_1 y_2}{x_2 - x_1}$$




RANSAC

1. selects m data items at random
2. estimates parameters $X=[a,b,...]$
3. finds how many data items K out of M fit the model with parameter vector X within a user given tolerance
4. if K is big enough, accept fit and exit with success
5. repeat 1..4 L times
6. failed, if you got here

$$d(ax + by + c = 0, (x_0, y_0)) = \frac{|ax_0 + by_0 + c|}{\sqrt{a^2 + b^2}}$$

$$d(P_1, P_2, (x_0, y_0)) = \frac{|(y_2 - y_1)x_0 - (x_2 - x_1)y_0 + x_2y_1 - y_2x_1|}{\sqrt{(y_2 - y_1)^2 + (x_2 - x_1)^2}}$$



T_T = tolerance threshold maximum distance between a point and the line

If $d < T_T \Rightarrow$ point fits the model

$Q > T_Q$ = stopping criteria e.g. 0.5


$Q = K / M$ = quality measure

K - number of inliers (or outliers)

RANSAC




RANSAC




- Select sample of m points at random

RANSAC




- Select sample of m points at random
- Calculate model parameters that fit the data in the sample

RANSAC




- Select sample of m points at random
- Calculate model parameters that fit the data in the sample
- Calculate error function d for each data point

RANSAC




- Select sample of m points at random
- Calculate model parameters that fit the data in the sample
- Calculate error function for each data point
- Select data that support current hypothesis, estimate Q

RANSAC




- Select sample of m points at random
- Calculate model parameters that fit the data in the sample
- Calculate error function for each data point
- Select data that support current hypothesis, estimate Q
- **Repeat sampling**

RANSAC



- Select sample of m points at random
- Calculate model parameters that fit the data in the sample
- Calculate error function for each data point
- Select data that support current hypothesis, estimate Q
- **Repeat sampling**

RANSAC



ALL-INLIER SAMPLE

RANSAC time complexity

$$O(m k m_s t_m M)$$

m ... number of data items to estimate model parameters

k ... number of iterations

t_m ... time to compute a single model

m_s ... average number of models per sample

M ... total number of items

RANSAC

- How big Q has to be depends on what percentage of the data you think belongs to the structure being fit and how many structures you have in the data.
 - If there are multiple structures then, after a successful fit, remove the fit data and redo RANSAC.

- You can find the number of iterations L by:

p_{fail} = probability of L consecutive failures

- $p_{fail} = (\text{prob that a given trial is a failure})^L$

- $p_{fail} = (1 - \text{prob that a given trial is a success})^L$

- $p_1 = \text{prob that a random data item fits the model}$ $p_1 = \left(\frac{\#inliers}{M} \right)$


- $p_g = \text{prob that a given trial is a success}$ $p_g = (p_1)^m$

- $p_{fail} = (1 - (\text{prob that a random data item fits the model})^m)^L$ $p_{fail} = (1 - (p_1)^m)^L$

$$L = \frac{\log(p_{preset\ fail})}{\log(1 - (p_1)^m)}$$

$$p_{preset\ fail} = 1 - p_{success}$$

Usually set very low e.g. 0.01



Hough Transform

Hough Transform


- Image with line structures
 - Point coordinates $P(x,y)$ given by white pixels



Hough Transform

- A method to find model parameters that fits the data
 - It was introduced in 1962 (Hough 1962) and first used to find lines in images a decade later (Duda 1972).



$$y=ax+b$$



$$(x-a)^2+(y-b)^2=r^2$$

Hough Transform




- Origin: Detection of straight lines in clutter
 - Basic idea: each candidate point votes for all lines that it is consistent with.
 - Votes are accumulated in quantized array
 - Local maxima correspond to candidate lines
- Representation of a line
 - Usual form $y = ax + b \Rightarrow b = y - ax$



Hough Transform


Representation of a line

- Usual form $y = ax + b \Rightarrow b = y - ax$ has a singularity around 90°
- Better parameterization: $x \cos(\theta) + y \sin(\theta) = \rho$




Hough Transform

- Each curve in the figure represents the family of lines that pass through a particular point (x_i, y_i) in the xy -plane.
- The intersection point (ρ', θ') corresponds to the lines that passes through two points (x_i, y_i) and (x_j, y_j)
- A horizontal line will have $\theta=0$ and ρ equal to the intercept with the y -axis.
- A vertical line will have $\theta=90$ and ρ equal to the intercept with the x -axis.




Hough Transform

- Noisy Line



Tokens




Votes

θ


- Problem: Finding the true maximum

Hough Transform

- Noise Input



Tokens




Votes

- Problem: Lots of spurious maxima

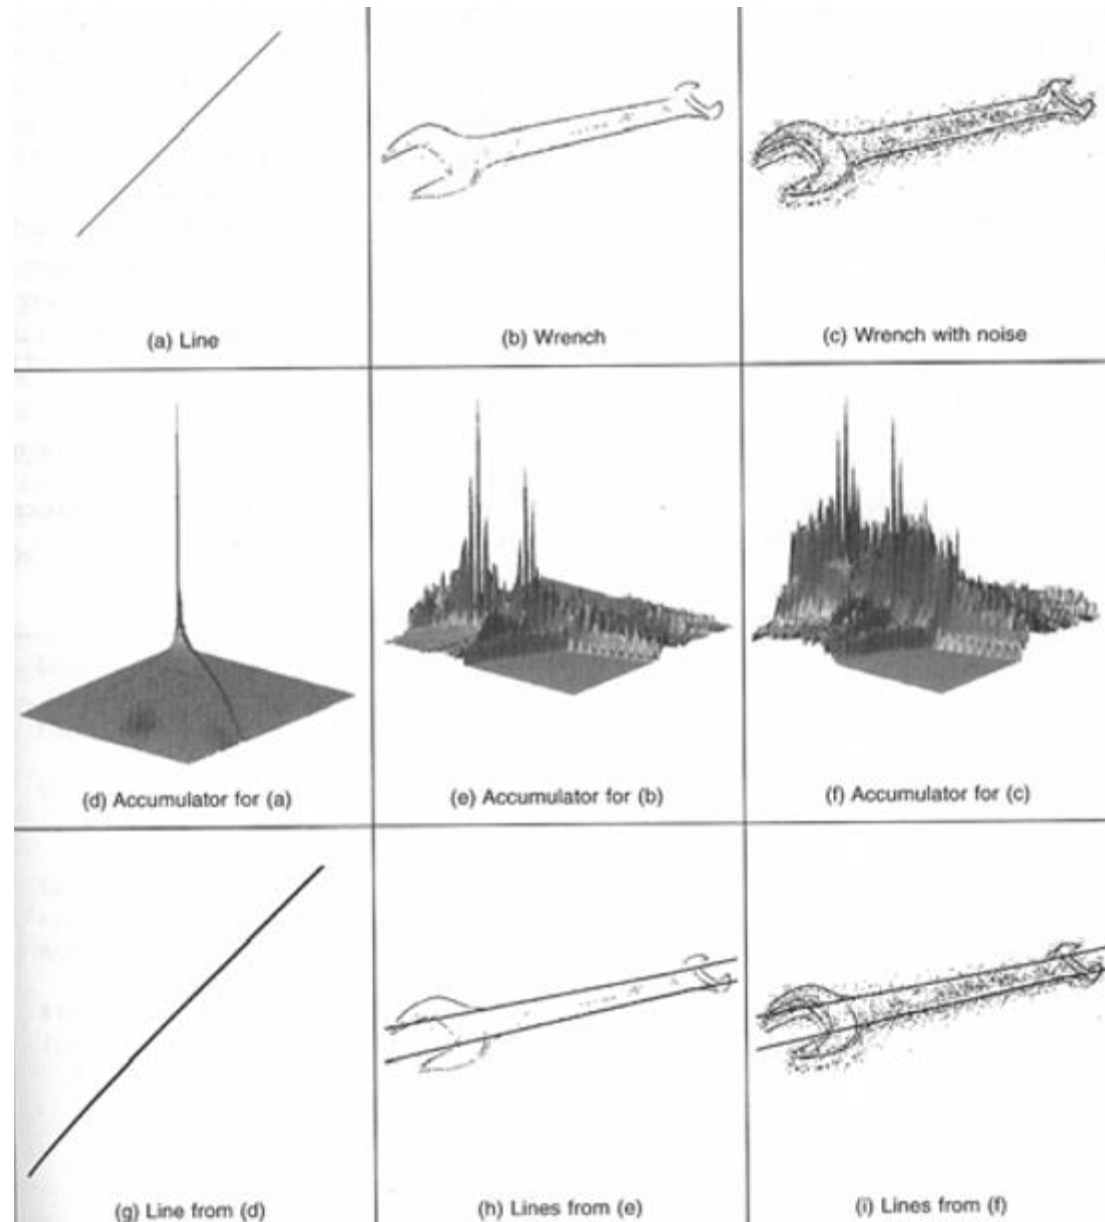
Hough Transform

- A square (left) and a circle (right)



Hough Transform

- Example images




Hough Transform

- Given data points find model parameters to fit to the data
 1. Define model and its parameters
 2. Derive a formula to obtain parameter values given a data point
 3. Define quantization steps and limits of the parameter space
 4. Build accumulative array for the parameter space (Hough space)
 5. Compute parameter values for a data sample and increment the cell in the array that corresponds to these values
 6. Repeat point 5 for every data sample
 7. Perform non maxima suppression in the array
 8. Recover the parameter values that correspond to the maxima.

Generalized Hough Transform

- Finding free form shapes which cannot be parameterised with a small number of parameters
 1. A reference point is chosen inside
 2. A line is constructed joining the reference point and the boundary
 3. The boundary direction is found at the point of intersection of line and boundary
 4. A reference table is constructed

ϕ	(r_1, α_1)		
ϕ_1	(r_1, α_1)	...	(r_1, α_1)
:	:	...	:
ϕ_k	(r_k, α_k)	...	(r_k, α_k)




Generalized Hough Transform

- Find location of the shape in the image
 - Construct 2D accumulative array for x,y location of reference point
 - Determine edge orientation at every edge point
 - Given one point use reference table to get corresponding (r, α) values
 - Calculate position of the reference point


$$x_r = x_\phi + r_\phi \cos(\alpha_\phi) \quad y_r = y_\phi + r_\phi \sin(\alpha_\phi)$$

ϕ	(r_1, α_1)		
ϕ_1	(r_1, α_1)	...	(r_1, α_1)
:	:	...	:
ϕ_k	(r_k, α_k)	...	(r_k, α_k)



Generalized Hough Transform

- Generalization for an arbitrary contour or shape
 - Choose reference point for the contour (e.g. center)
 - For each point on the contour remember where it is located w.r.t. to the reference point
 - E.g. if the center is the reference point: remember radius r and angle relative to the tangent of the contour
 - Recognition: whenever you find a contour point, calculate the tangent angle and 'vote' for all possible reference points
 - Instead of reference point, can also vote for transformation
- ⇒ The same idea can be used with local features!



When is the Hough transform useful?

- Textbooks wrongly imply that it is useful mostly for finding lines
 - In fact, it can be very effective for recognizing arbitrary shapes or objects
- The key to efficiency is to have each feature (token) determine as many parameters as possible
 - For example, lines can be detected much more efficiently from small edge elements (or points with local gradients) than from just points
 - For object recognition, each token should predict location, scale, and orientation (4D array)
- Bottom line: The Hough transform can extract feature groupings from clutter in linear time!

Comparison

Gen. Hough Transform

More robust.

Advantages

- Can be very effective for recognizing arbitrary patterns
 - Can handle high percentage of outliers (>95%)
 - Extracts groupings from clutter in linear time
- ### Disadvantages
- Quantization issues
 - Only practical for small number of dimensions (up to 4)
 - Handles missing and occluded data very gracefully.
 - Can be adapted to many types of forms, not just lines

Improvements available

- Probabilistic Extension
 - Continuous Voting Space
- } [Leibe04]

RANSAC

Advantages

- General method suited to large range of problems
- Conceptually simple and easy to implement
- Independent of number of dimensions

Disadvantages

- Only handles moderate number of outliers (<50%)
- Many variants available, e.g.
 - PROSAC: Progressive RANSAC [Chum05]
 - Preemptive RANSAC [Nister05]

Applications

- Sony Aibo
(Evolution Robotics)
 - Recognize docking station
 - Communicate with visual cards

AIBO® Entertainment Robot

Official U.S. Resources and Online Destinations

ERS-7
Entertainment Robot AIBO




ERS-7 with:
Wireless LAN
AIBO MIND software
Energy Station
AIBOne
Pink Ball
AIBO Cards (15)
WLAN Manager CD
Battery & AC Adapter

3rd Generation
Pre-order Now!

Example: Finding Feature Matches

- Find best stereo match within a square search window (here 300 pixels^2)



from Hartley & Zisserman


Example: Finding Feature Matches

- Find best stereo match within a square search window (here 300 pixels^2)

before RANSAC



after RANSAC



from Hartley & Zisserman